

Lexical Analysis

Definition: **Lexical Analysis** is the operation of dividing the input program into sequence of lexemes(tokens)

Input

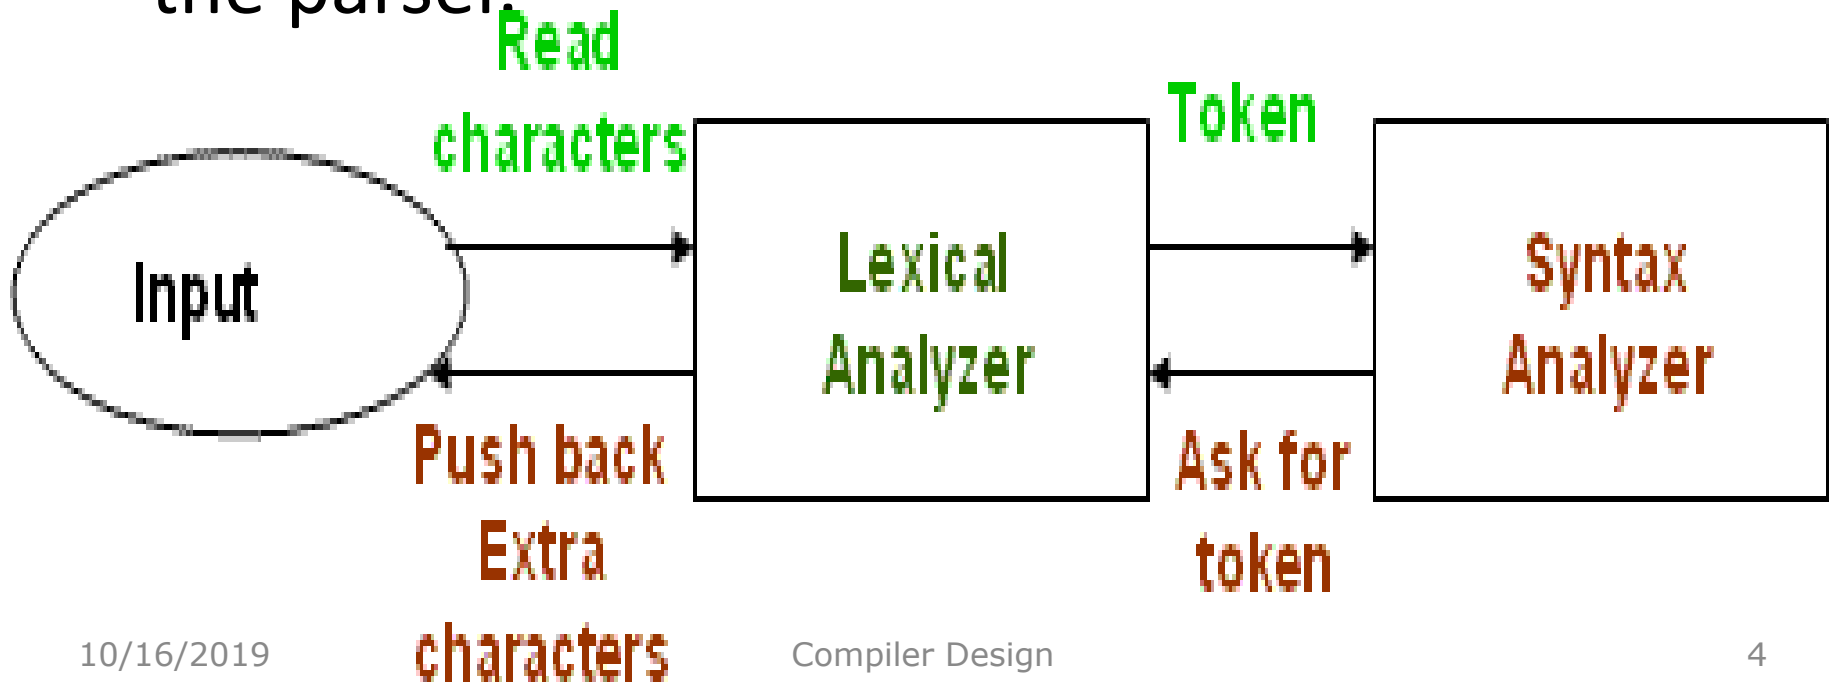
- Read string input
 - Might be sequence of characters
 - Might be sequence of lines
 - Character set:

The Output

- A series of *tokens*: *kind, location, name* (if any)
 - Punctuation () ; , []
 - Operators + - ** :=
 - Keywords begin end if while try catch
 - Identifiers Square_Root
 - String literals “press Enter to continue”
 - Character literals ‘x’
 - Numeric literals
 - Integer: 123
 - Floating_point: 4_5.23e+2
 - Based representation: 16#ac#

Introduction

- The role of the **lexical analyzer** is to read a sequence of characters from the source program and **produce tokens** to be used by the parser.



- The stream of tokens is sent to the parser for syntax analysis
- The lexical analyzer also interacts with the symbol table, e.g., when the lexical analyzer discovers a lexeme constituting an identifier, it needs to enter that lexeme into the symbol table

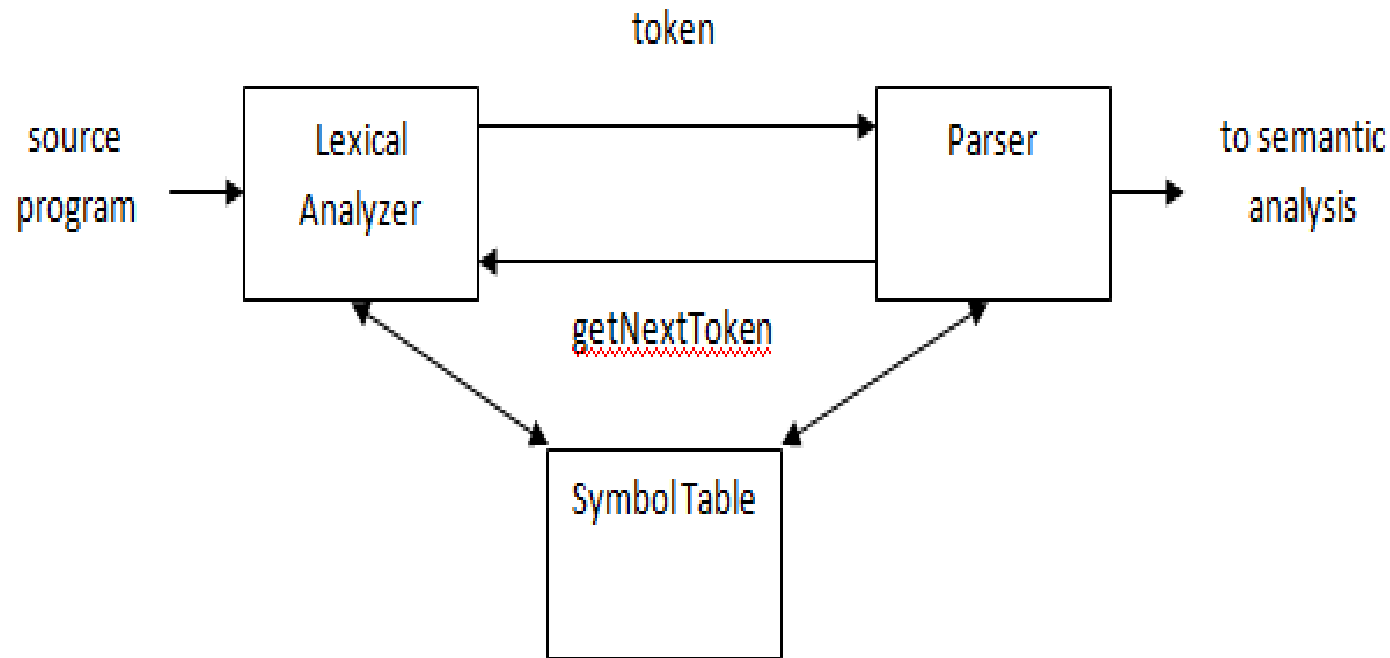


Figure 2.1: Interactions between the lexical analyzer and the parser

- “getNextToken” is a command sent from the parser to the lexical analyzer.
- On receipt of the command the lexical analyzer scans the input until it determines the next token, and returns it.
- It skips comments and whitespaces while creating these tokens

Introduction (cont'd)

- The lexical analyzer can also perform the following secondary tasks:
 - stripping out blanks, tabs, new lines
 - stripping out comments
 - keeping track of line numbers **(to correlate an error with the source and line number)**
 - Expanding macros in some lexical analyzers

Token, pattern, lexeme

- **Token**: a sequence of characters from the source program having a **collective meaning**.
- A classification for a common set of strings
- Examples Include <Identifier>, <number>, etc
- A single token can be produced by different sequences of characters.
- E.g. x, distance, count → IDENT
- Tokens are terminals in the grammar of the source language.

Token, pattern, lexeme (cont'd)

- **Pattern:** a rule describing the set of lexemes that represent a token.
- Patterns are usually specified using regular expressions.

Example

`[a-zA-Z]*`

Lexemes matched: a, ab, count, ...

Token, pattern, lexeme (cont'd)

- **Lexeme**: a sequence of characters in the source program that is matched by a pattern for a certain token.
- Lexeme: The smallest logical units(words) of a program
 - eg: i, sum, for,10,++, "%d\n",<=.

Example:

x, distance, count → IDENT

token: IDENT

lexemes: x, distance, count

Token, pattern, lexeme (cont'd)

- Example: Some tokens and their lexemes in Pascal (a high level, case insensitive programming language)

Token	Some lexemes	Informal pattern
begin	begin, Begin, BEGIN, beGin, ...	Begin in small or capital letters
if	if, IF, iF, If	if in small or capital letters
ident	Distance, F1, x, Dist1, ...	Letter followed by zero or more letters and/or digits

Complete the above table for the C++ programming language.

Token, pattern, lexeme (cont'd)

- In general, in programming languages, the following are tokens:
 - keywords
 - operators
 - identifiers
 - constants
 - literals
 - punctuation symbols
 - ...

Attributes of tokens

- When more than one lexeme matches a pattern, the scanner must provide additional information about the particular lexeme to the subsequent phases of the compiler.
- For ex., both **0** and **1** match the pattern for the token **num**. But the code generator needs to know which number is recognized.

Attributes of tokens (cont'd)

- The lexical analyzer collects information about tokens into their associated attributes.
- Practically, a token has one attribute: a pointer to the symbol table entry in which the information about the token is kept.
- **Symbol table entry** contains information about the token such as the lexeme, the line number in which it was first seen, ...

Attributes of tokens (cont'd)

- For ex. consider $x = y + 2$

The tokens and their attributes are written as:

<**id**, pointer to symbol-table entry for x>

<**assign_op**, >

<**id**, pointer to symbol-table entry for y>

<**plus_op**, >

<**num**, integer value 2>

- **Example:** $E = M * C ** 2$

<id, pointer to symbol-table entry for E>

<assign_op, >

<id, pointer to symbol-table entry for M>

<mult_op, >

<id, pointer to symbol-table entry for C>

<exp_op, >

<num, integer value 2>

Errors

- Very few errors are detected by the lexical analyzer.
- For ex., if the programmer mistakes **wihle** for **while**, the lexical analyzer cannot detect the error (why?)
- Nonetheless, if a certain sequence of characters follows none of the specified patterns, the lexical analyzer can detect the error.
- **Examples of Lexical Errors:**
 - **Illegal characters in** source program
 - For ex. if there is a ? symbol in the source program and no pattern contains ?
 - **Exceeding length of identifier or numeric constants**
 - **Unterminated strings or comments**

Handling Lexical Errors

- When an error occurs, the lexical analyzer recovers by:
 - **Panic mode recovery** :skipping (deleting) successive characters from the remaining input until the lexical analyzer can find a well-formed token (panic mode recovery)
 - deleting extraneous characters
 - inserting missing characters
 - replacing an incorrect character by a correct character
 - transposing two adjacent characters

Example – tokens in Java

1. Identifier: A *Javaletter* followed by zero or more *Javaletterordigits*.
A *Javaletter* includes the characters a-z, A-Z, _ and \$.
2. Constants:
 - 2.1 Integer Constants
 - ▶ Octal, Hex and Decimal
 - ▶ 4 byte and 8 byte representation
 - 2.2 Floating point constants
 - ▶ float - ends with f
 - ▶ double
 - 2.3 Boolean constants – true and false
 - 2.4 Character constants – 'a', '\u0034', '\t'
 - 2.5 String constants – "", "\"", "A string".
 - 2.6 Null constant – null.
3. Delimiters: (,), {, }, [,] , ;, . and ,
4. Operators: =, >, < ...>>>=
5. Keywords: abstract, boolean ... volatile, while.



Recap - Token Attributes

Apart from the token itself, the lexical analyser also passes other informations regarding the token. These items of information are called *token attributes*

EXAMPLE

lexeme	<token, token attribute>
3	< const, 3>
A	<identifier, A>
if	<if, ->
=	<assignop, ->
>	<gt, ->
;	<semicolon, ->

Specifying and recognizing tokens

- Regular expressions are used to specify the patterns of tokens.
- **Regular Definitions**
- For notational convenience, we may wish to give names to certain regular expressions and use those names in subsequent expressions, as if the names were themselves symbols
- If Σ is an alphabet of basic symbols, then a *regular definition* is a sequence of definitions of the form:

$$d_1 \rightarrow r_1$$

$$d_2 \rightarrow r_2$$

...

$$d_n \rightarrow r_n$$

where:

- Each d_i is a new symbol, not in Σ and not the same as any other of the d 's, and
- Each r_i is a regular expression over the alphabet $\Sigma \cup \{d_1, d_2, \dots, d_{i-1}\}$

- **Example:**

C identifiers are strings of letters, digits, and underscores. Here is a regular definition:

letter_ $\rightarrow A \mid B \mid \dots \mid Z \mid a \mid b \mid \dots \mid z \mid _$

digit $\rightarrow 0 \mid 1 \mid \dots \mid 9$

id $\rightarrow letter_ (letter_ \mid digit)^*$

- Regular expressions (RE) defined by an alphabet (terminal symbols) and three operations:

- Alternation $RE_1 \mid RE_2$

- Concatenation $RE_1 RE_2$

- Repetition RE^* (zero or more RE's)

- Extensions of Regular Expressions
- **One or more instances.** The unary postfix operator $+$ represents the positive closure of a regular expression and its language. The operator $+$ has the same precedence and associativity as the operator $*$. $r^* = r^+ \mid \lambda$ and $r^+ = rr^* = r^*r$
- **Zero or one instance.** The unary postfix operator $?$, means “zero or one occurrence”. That is $r?$ is equivalent to $r \mid \lambda$. The $?$ operator has the same precedence and associativity as $*$ and $+$
- **Character classes.** A regular expression $a_1 \mid a_2 \mid \dots \mid a_n$, where a_i 's are each symbols of the alphabet, can be replaced by the shorthand $[a_1a_2 \dots a_n]$. More importantly, when a_1, a_2, \dots, a_n form a logical sequence, e.g., consecutive uppercase letter, lowercase letters, or digits, we can replace them by a_1 - a_n , that is, just the first and the last separated by hyphen. Thus $[abc]$ is shorthand for $a \mid b \mid c$, and $[a-z]$ is shorthand for $a \mid b \mid \dots \mid z$

Specifying RE's in Unix Tools

- Single characters `a b c d \x`
- Alternation `[bcd] [b-z] ab|cd`
- Any character `.` *(period)*
- Match sequence of characters `x* y+`
- Concatenation `abc[d-q]`
- Optional RE `[0-9]+(\.[0-9]*)?`

Example

letter $\rightarrow A|B|C|\dots|Z|a|b|c|\dots|z$

digit $\rightarrow 0|1|\dots|9$

identifier $\rightarrow \text{letter} (\text{letter}|\text{digit})^*$

Specifying (cont'd)

- Lex provides shortcuts for describing regular expressions in a compact manner.

Example

[a-z] stands for $a|b|c|\dots|z$

[0-9] stands for $0|1|\dots|9$

[abc] stands for $a|b|c$

Creating a Lexical Analyzer

Two approaches:

1. *Hand code* – This is only of historical interest now.
 - ▶ Possibly more efficient.
2. *Use a generator* – To generate the lexical analyser from a formal description.
 - ▶ The generation process is faster.
 - ▶ Less prone to errors.

Automatic Generation of Lexical Analysers

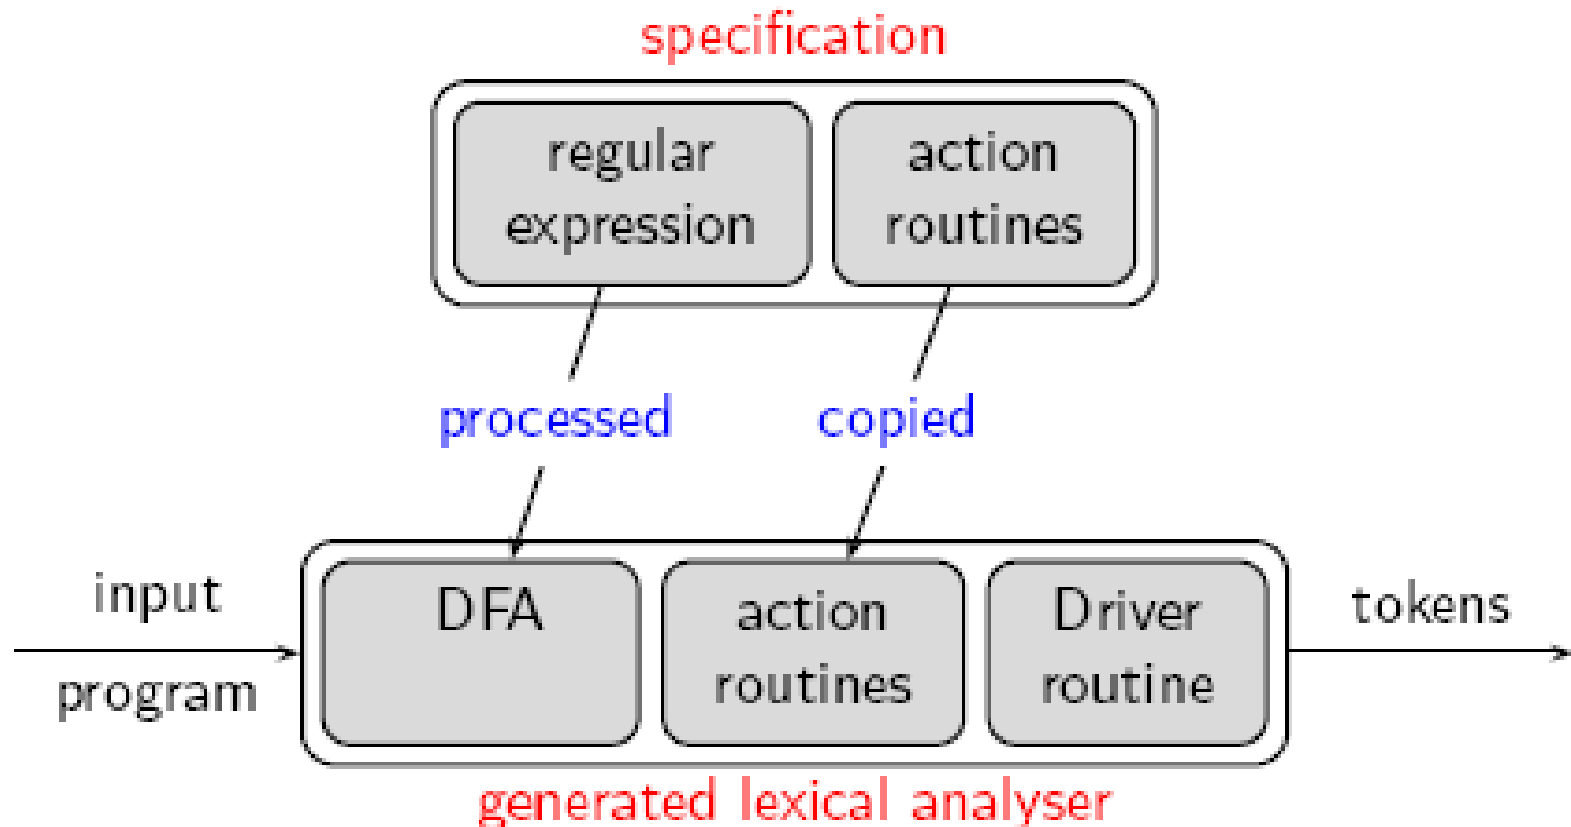
Inputs to the lexical analyser generator:

- A **specification** of the tokens of the source language, consisting of:
 - ▶ a **regular expression** describing each **token**, and
 - ▶ a **code fragment** describing the **action to be performed**, on identifying each token.

The generated lexical analyser consists of:

- A **deterministic finite automaton (DFA)** constructed from the token specification.
- A **code fragment** (a driver routine) which can traverse *any DFA*.
- Code for the **action specifications**.

Automatic Generation of Lexical Analysers




Example of Lexical Analyser Generation

Suppose a language has two tokens

Pattern	Action
<code>a*b</code>	<code>{ printf("Token 1 found"); }</code>
<code>c+</code>	<code>{ printf("Token 2 found"); }</code>


From the description, construct a structure called a deterministic finite automaton (DFA).



Example of Lexical Analyser Generation

Now consider the following together:

The DFA



The driver routine

```
void nexttoken ()
{state = 0; c = nextchar();
 while (valid(nextstate[state,c]))
 {state = nextstate[state,c];
  c = nextchar();}
 if (!final(state))
 {error; return;}
 else
 {unput(c);action();return;}}
```

The actions

```
void action();
{
 switch(state)
  2:{printf("Token 1 found");
    break;}
  3:{printf("Token 2 found");
    break;}
}
```

The input and output


Input: aabadbcc ←

Output:

Example of Lexical Analyser Generation

Now consider the following together:

The DFA



The driver routine

```
void nexttoken ()
{state = 0; c = nextchar();
 while (valid(nextstate[state,c]))
 {state = nextstate[state,c];
  c = nextchar();}
 if (!final(state))
 {error; return;}
 else
 {unput(c);action();return;}}
```


The actions

```
void action();
{
 switch(state)
 2:{printf("Token 1 found");
   break;}
 3:{printf("Token 2 found");
   break;}
}
```


The input and output

Input: aabadbcc←

Output: Token 1 found
Token 1 found
Token 2 found



- Fig:Transition diagram for relop



- Fig:diagram for unsigned numbers in Pascal

Exercise

- All Strings that start with “ab” and end with “ba”
- All Strings in Which {1,2,3} exist in ascending order