

Chapter Two

Transaction

What is a Transaction?

- **A Transaction:**
 - Logical unit of database processing that includes one or more access operations (read -retrieval, write - insert or update, delete).
- A transaction (set of operations) may be stand-alone specified in a high level language like SQL submitted interactively, or may be embedded within an application program.
- **Transaction boundaries:**
 - Any single transaction in an application program is bounded with **Begin** and **End** statements.
- An **application program** may contain several transactions separated by the Begin and End transaction boundaries

Transaction operations

- the basic database access operations that a transaction can include are as follows:
 - `read_item(X)`.
 - ✓ Reads a database item named X into a program variable
 - `write_item(X)`.
 - ✓ Writes the value of program variable X into the database item named X.
- The basic access is one disk block from disk to memory

read_item(*X*)

- Executing a read_item(*X*) *command includes* the following steps:
 1. Find the address of the disk block that contains *x*
 2. Copy the disk block to memory buffer(if not in memory)
 3. Copy the item from the buffer to program variable *x*

write_item(X)

- Executing a `write_item(X)` *command includes the following steps:*
 1. Find the address of the disk block that contains x
 2. Copy the disk block to memory buffer (if not in memory)
 3. Copy the program variable x into buffer
 4. Store the updated buffer to disk

A Transaction: An Informal Example

- E.g. Transfer transaction to transfer 500 birr from account A to account B: For a user it is one activity
- To database:
 - Read balance of A : read(A)
 - Read balance of B : read (B)
 - Subtract 500 birr from A
 - Add 500 birr to B
 - Write new value of A back to disk
 - Write new value of B back to disk

Cont'd

- Transaction to transfer 500 birr from account A to account B:
 - 1.**read**(A)
 2. $A := A - 500$
 - 3.**write**(A)
 - 4.**read**(B)
 5. $B := B + 500$
 - 6.**write**(B)

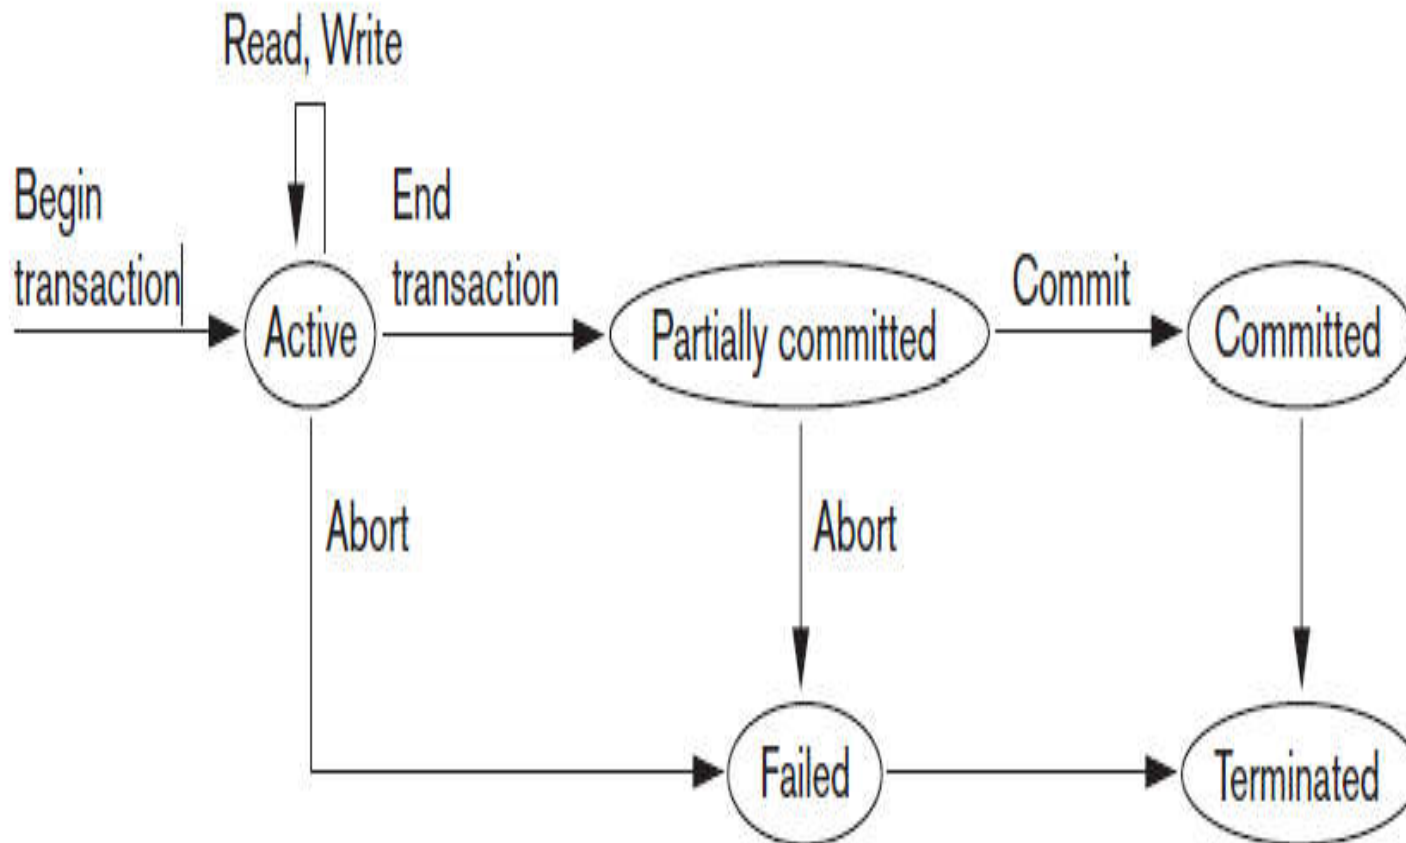
Transaction states and additional operations

- For recovery purposes, it keeps track of start, terminate and commits or aborts
- the recovery manager of the DBMS needs to keep track of the following operations :
 - ✓ BEGIN_TRANSACTION
 - ✓ READ OR WRITE
 - ✓ END_TXN
 - ✓ COMMIT_TXN
 - ✓ ROLLBACK or ABORT
- Failed state need rollback operation

Transaction States

- transaction must be in one of several states:
 - **Active** : A transaction that has executed at least one action, but has not completed.
 - **Partially committed** : A transaction that has completed all actions, but has not been confirmed as completed.
 - **Failed** : A transaction that fails completion checks, or is aborted during the active state.
 - **Committed** : A completed and checked transaction
 - **Aborted** : A transaction that has failed and has had no effect on the database system

Cont'd



Desirable Properties of Transactions

- Transactions should possess several properties, often called the **ACID properties**;
- The following are the ACID properties
 - ✓ **Atomicity**
 - ✓ **Consistency preservation**
 - ✓ **Isolation**
 - ✓ **Durability or permanency**

Cont'd

- **Atomicity**

- ✓ A transaction is an atomic unit of processing;
- ✓ it should either be performed in its entirety or not performed at all

- **Consistency preservation.**

- ✓ A transaction should be consistency preserving,
- ✓ it is completely executed from beginning to end without interference from other transactions,

Cont'd

- **Isolation**
 - ✓ A transaction should appear as though it is being executed in isolation from other transactions,
 - ✓ even though many transactions are executing concurrently
- **Durability or permanency.**
 - ✓ The changes applied to the database by a committed transaction must persist in the database.
 - ✓ These changes must not be lost because of any failure.

Schedules

- Multiple transactions can be executed concurrently by interleaving their operations
- **Schedule S** – a sequences of instructions that specify the **chronological order** in which instructions of concurrent transactions T_1, T_2, \dots, T_n are executed
 - A schedule for a set of transactions **must consist** of all instructions of those transactions.
 - Must **preserve** the order in which the instructions appear in each individual transaction.

Cont'd

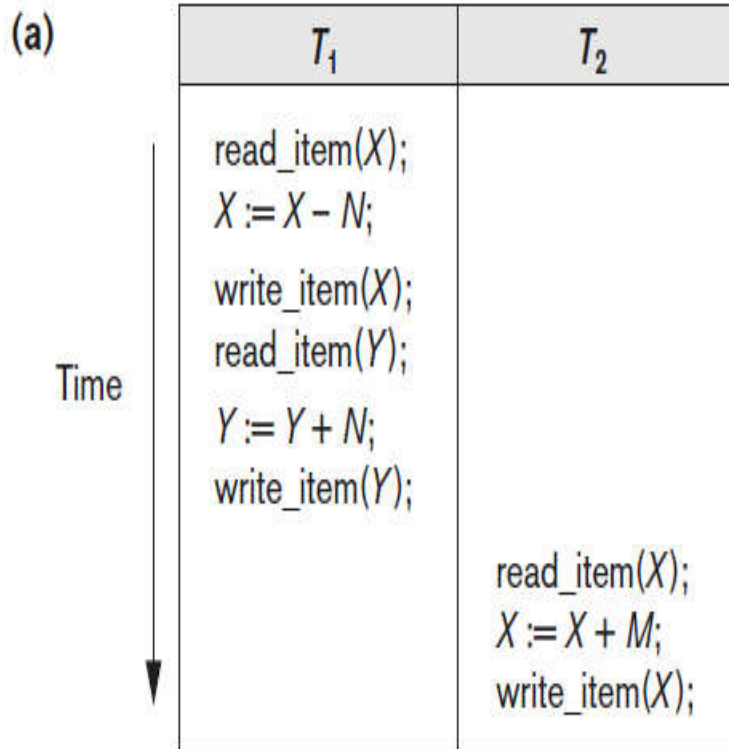
Serial schedule

- ✓ S is a schedule that executes all of the operations from one transaction T1, before moving on to the operations of another transaction T2.

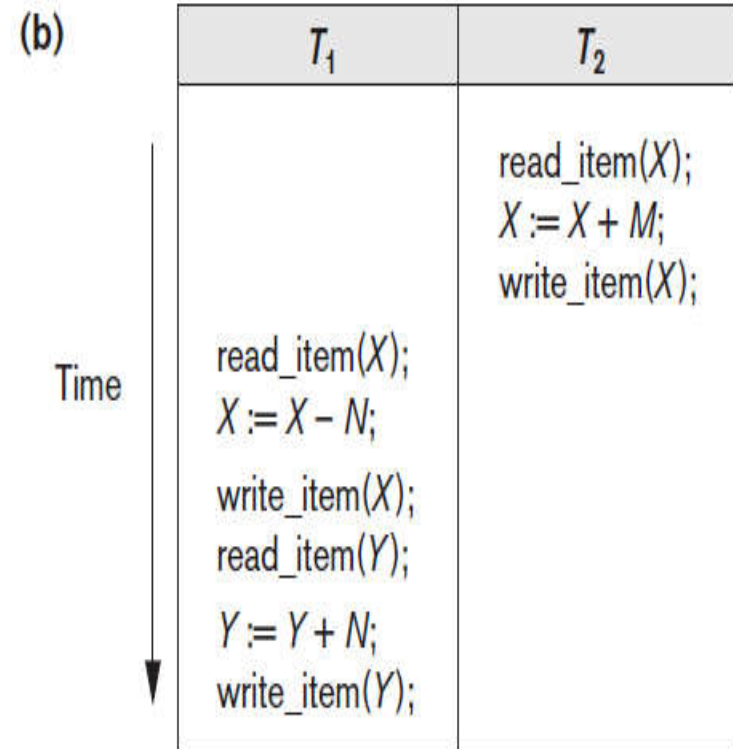
Non-serial schedule(*interleaved* schedule)

- ✓ An is a schedule in which the operations of an individual transaction are executed in order with respect to the same transaction,
- ✓ This is a type of Scheduling where the operations of multiple transactions are interleaved.
- ✓ This might lead to a rise in the concurrency problem

Serial schedule

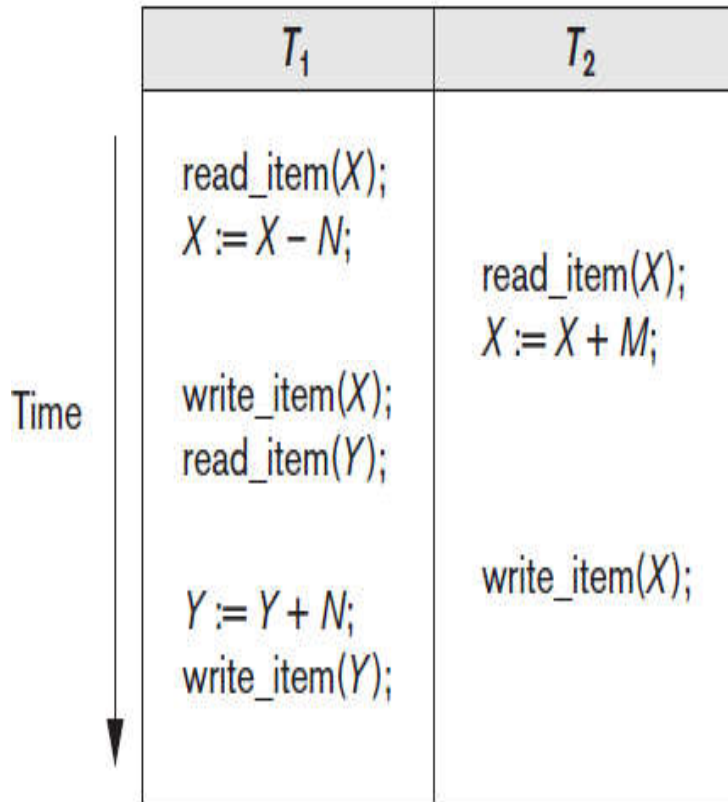


Schedule A

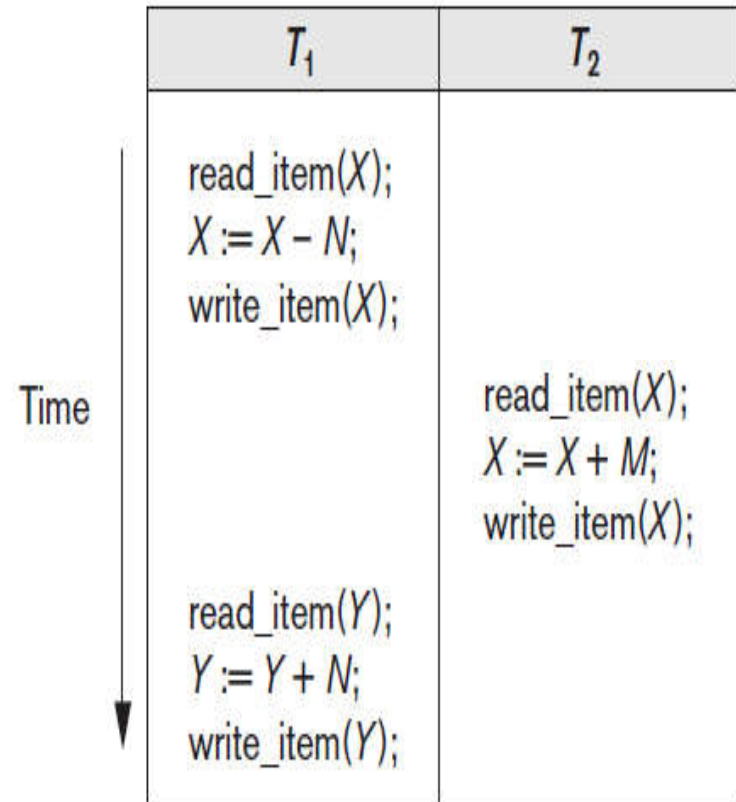


Schedule B

Non-serial schedules



Schedule C



Schedule D

Result Equivalent Schedules

- Two schedules are result equivalent if they produce the same final state of the database
- Problem: May produce same result by accident!

S1

```
read_item(X);  
X:=X+10;  
write_item(X);
```

S2

```
read_item(X);  
X:=X*1.1;  
write_item(X);
```

Schedules S1 and S2 are result equivalent for $X=100$ but not in general

Characterizing Schedules based on Serializability

- The concept of Serializable of schedule is used to identify which schedules are correct when concurrent transactions executions have interleaving of their operations in the schedule
- Serial schedule:
 - A schedule S is serial if, for every transaction T participating in the schedule, all the operations of T are executed consecutively in the schedule.
 - Otherwise, the schedule is called nonserial schedule.
- Serializable schedule:
 - a schedule whose effect on any consistent database instance is identical to that of some complete **serial schedule** over the set of *committed* transactions in S .
 - A nonserial schedule **S is serializable** is equivalent to say that it is **correct** to the result of one of the serial schedule

Serializability of Schedules

Serializability

- A schedule S of n transactions is **serializable** if it is *equivalent to some serial schedule* of the same n transactions
- **Kinds of serializability**
 - Conflict-serializable schedule
 - View-serializable schedule

Result Equivalent Schedules

- Two schedules are result equivalent if they produce the same final state of the database
- Problem: May produce same result by accident!

S1

```
read_item(X);  
X:=X+10;  
write_item(X);
```

S2

```
read_item(X);  
X:=X*1.1;  
write_item(X);
```

Schedules S1 and S2 are result equivalent for $X=100$ but not in general

Two types of equivalent schedule:

- Two schedules are said to be conflict equivalent if the order of any two conflicting operations is the same in both schedules.
- Eg
 - $S1: r1(x); w2(x)$ & $S2: w2(x); r1(x)$
 - $S1: w1(x); w2(x);$ & $S2: w2(x); w1(x);$
- Conflict serializable:
 - A schedule S is said to be conflict serializable if it is **conflict equivalent** to some serial schedule S' .
 - Every **conflict serializable** schedule is **serializable**.

Not conflict
equivalent

Conflict Serializable

- Schedule S is conflict serializable if it is conflict equivalent to some *serial* schedule S'
 - We can reorder the *non-conflicting* operations to improve efficiency
- Non-conflicting operations:
 - Reads and writes from same transaction
 - Reads from different transactions
 - Reads and writes from different transactions on different data items
- Conflicting operations:
 - Reads and writes from different transactions on same data item

- Two operations in a schedule are said to be **conflict** if they satisfy all the three of the following conditions.

- They belong to different transactions
- They access the same data item X
- At least one of the operations is a write_Item(X)

Eg. **S_a**: r1(X); r2(x); w1(X); r1(Y); W2(X); W1(Y);

- r1(X) and w2(X)
 - r2(X) and w1(X);
 - W1(X) and w2(X)
- } **Conflict operations**

- r1(X) and r2(X)
 - W2(X) and w1(Y)
- } **No Conflict**

Cont'd

Testing for conflict serializability: Algorithm

- Looks at only read_Item (X) & write_Item (X) operations
- Constructs a precedence graph (serialization graph) - a graph with directed edges
- An edge is created from T_i to T_j if one of the operations in T_i appears before a conflicting operation in T_j
- The schedule is serializable if and only if the precedence graph has no cycles.

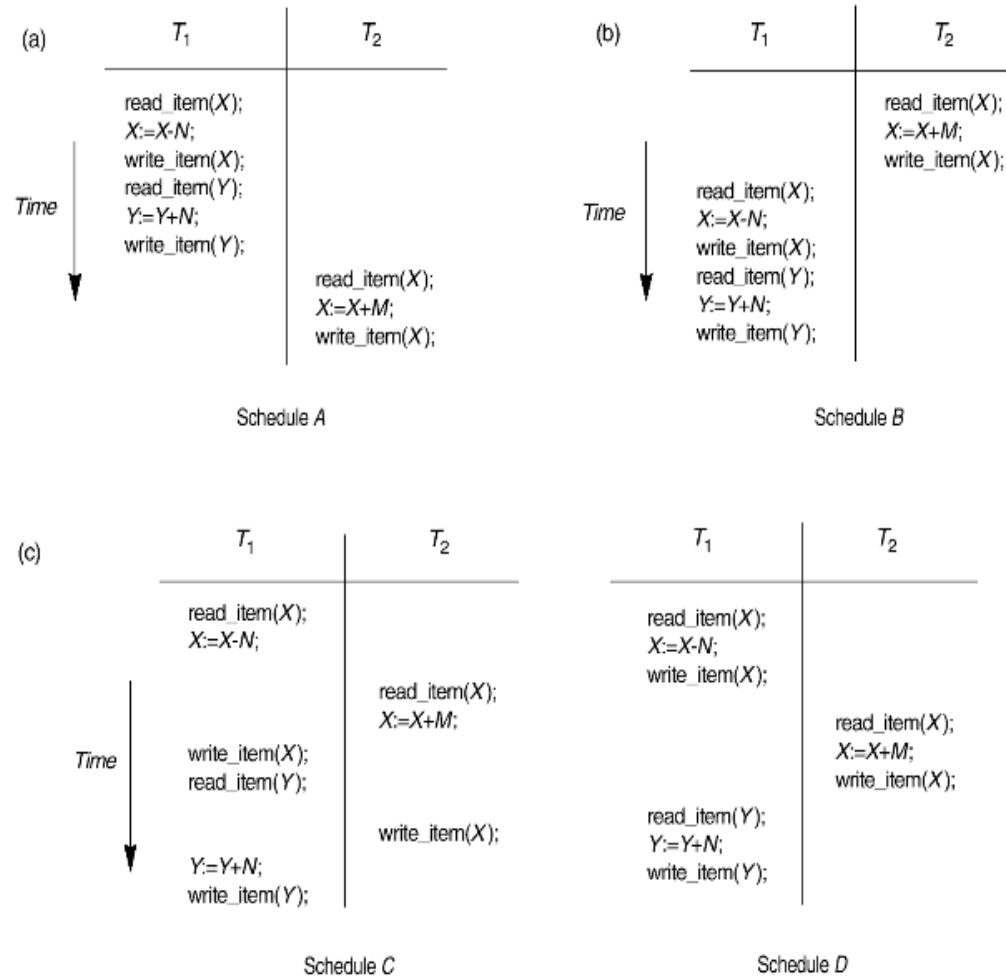






Fig 1:

Examples of serial and nonserial schedules involving transactions T_1 and T_2 .

Constructing the Precedence Graphs

- FIGURE 2: Constructing the precedence graphs for schedules A and D from Figure 1 to test for conflict serializability.
 - (a) Precedence graph for serial schedule A.
 - (b) Precedence graph for serial schedule B.
 - (c) Precedence graph for schedule C (not serializable).
 - (d) Precedence graph for schedule D (serializable, equivalent to schedule A).




View Equivalence and View Serializability

- **View Serializability** is a process to find out that a given **schedule** is view serializable or not.
 - To check whether a given schedule is view serializable,
 - we need to check whether the given schedule is **View Equivalent** to its serial schedule.
 - Definition of serializability based on view equivalence.
- A schedule is *view serializable* if it is *view equivalent* to a serial schedule
- A schedule is *view serializable* if it is *view equivalent* to a serial schedule

view equivalence:

- As long as each read operation of a transaction reads the result of *the same write operation* in both schedules,
- the write operations of each transaction must produce the same results.
- **“The view”**: the read operations are said to see the *same view* in both schedules.

- Two schedules are said to be view equivalent if the following three conditions hold:
 - The same set of transactions participates in S and S' , and S and S' include the same operations of those transactions.
 - If T_i reads a value A written by T_j in S_1 , it must also read the value of A written by T_j in S_2
 - for each data object A , the transaction that perform the final write on x in S_1 must also perform the final write on A in S_2



- **Relationship between view and conflict equivalence:**
 - The two are same under **constrained write assumption** which assumes that if T writes X, it is constrained by the value of X it read;
i.e., $\text{new } X = f(\text{old } X)$
 - Conflict serializability is **stricter** than view serializability.
 - a schedule that is view serializable is not necessarily conflict serializable.
 - Any conflict serializable schedule is also view serializable, but not vice versa.

Reading Assignment

Strict schedule