

Chapter 3

Supervised Learning

Supervised Learning

- Supervised learning: learning a function that maps an input to an output based on example **input-output** pairs.
- Supervised learning:
 - **Regression** problem, if the target variable is **continuous**
 - Regression algorithms are used to **predict the continuous** values such as price, salary, age, etc.
 - You can estimate the selling price of your house in your city.
 - **Classification** problem, if the target variable is categorical,
 - Classification algorithms are used to predict/Classify the discrete values such as Male or Female, True or False, Spam or Not Spam, etc.

Background – Terminology

- Let's review some common ML terms.
- Data is usually represented with a **feature matrix**.
 - **Features**
 - Attributes used for analysis
 - Represented by columns in feature matrix
 - **Instances/feature vector**
 - Entity with certain attribute values
 - Represented by rows in feature matrix
 - **Class labels**
 - Indicate category for each instance.
 - This example has two classes (C1 and C2).
 - Only used for supervised learning

		Features				
		Attributes used to classify instances				
		F1	F2	F3	F4	F5
Each instance has a class label	C1	41	1.2	2	1	3.6
	C2	63	1.5	4	0	3.5
	C1	109	0.4	6	1	2.4
	C1	34	0.2	1	0	3.0
	C1	33	0.9	6	1	5.3
	C2	565	4.3	10	0	3.2
	C1	21	4.3	1	0	1.2
	C2	35	5.6	2	0	9.1
Instances						

Fig. 1 Feature matrix

Classification

- The classification problem aims to predict group membership (i.e., labels or classes), for a set of observations.
- Classification is a two-step process: a **model construction (learning)** phase, and a **model usage (applying)** phase. Sometimes model adjustment/by validation considered as 3rd step.
- **Formally:**
Given a set of data-point $\mathbf{X} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ and a finite set of target classes $\mathbf{Y} = \{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_m\}$, Classification problems is to define a mapping $\mathbf{f} : \mathbf{X} \rightarrow \mathbf{Y}$, where each \mathbf{x}_i is assigned to one of class (\mathbf{y}_i).

Single label vs multi-label Classification

- The classification problem can be further divided into two categories:
 - Single-label classification : *only one label*
 - Multi-label classification: *more than one label*

Single-label data: $Y_1 \in \{0,1\}$

X_1	X_2	...	X_p	Y_1
0.12	1	...	12	0
2.34	9	...	-5	1
1.22	3	...	40	1
2.18	2	...	8	?

p input variables

Single output variable

Vs.

Multi-label Data : $[Y_1, \dots, Y_q] \in 2^q$

X_1	X_2	...	X_p	Y_1	Y_2	...	Y_q
0.12	1	...	12	0	1	...	1
2.34	9	...	-5	1	1	...	0
1.22	3	...	40	1	0	...	1
2.18	2	...	8	?	?	...	?

p input variables



q binary output variables

m training examples

Unknowns

Single label classification

- **Single label classification:** learning from a **set of data points** that are associated with only one target label from a set of disjoint labels Y , with $|Y| \geq 2$.
 - If $|Y| = 2$, the learning problem is called **binary classification**
e.g disease diagnosis, spam, malware detection, etc.
 - if $|Y| > 2$, then it is called a **multi-class classification** problem
e.g Language identification, an animal image classification, character recognition, face recognition, etc.




Multi-label classification

- What is the issue with Single-label classification?

Ans: In many real-world applications, each data instance can be associated with multiple class variables

- Examples:
 - A news article may cover multiple topics, such as politics, Agriculture, and economics

- Class1- Red Vs Blue
- Class2- Triangle Vs Oval



- So, how can we handle such multi-label problem?

Approaches to solve multi-label classification problem

- **Problem transformation** : Transforms multi-label classification problem to multiple single-label classification problems which can be called **Adapt data to the algorithm**.
 - **Binary Relevance (BR)**: Multi-label into single-label problems.
 - **Label Power-set (LP)** : Multi-label into multi-class – consider each label-sets as class.
 - **Classifier Chains (CC)**: resolves the BR limitations, by making label correlation task.

Other solutions [will not covered in this chapter]

- **Adapted Algorithm**: Perform multi-label classification, rather than transforming the problem into different subsets of problem.
- **Ensemble Approaches**: learning multiple classifier systems train multiple hypotheses to solve the same problem.

Binary Relevance

- Transform multi-label into single-label problems
- Suppose a training dataset D having five instances with an input feature vector X (x1 and x2) and output class vector Y (y1,y2,y3).

X1	X2	Y1	Y2	Y3
0.3	0.7	1	1	0
0.9	0.3	1	1	0
0.6	0.6	0	0	1
0.4	0.1	0	1	0
0.5	0.2	1	0	1

To multiple singl-label problems

- $X \rightarrow Y1$
- $X \rightarrow Y2$
- $X \rightarrow Y3$

Advantages

- Computationally efficient

Disadvantages

- Does not capture the dependence relations among the class variables

Label Powerset Method

- Transform each label combination to a class value and then learn a multi-class classifier with the new class values

X1	X2	Y1	Y2	Y3
0.3	0.7	1	1	0
0.9	0.3	1	1	0
0.6	0.6	0	0	1
0.4	0.1	0	1	0
0.5	0.2	1	0	1

Combine label to a class value

X1	X2	Y _{comb}
0.3	0.7	1
0.9	0.3	1
0.6	0.6	2
0.4	0.1	3
0.5	0.2	4

• $X \rightarrow Y_{\text{comb}}$

Advantages

- Learns the full joint of the class variables and each of the new class values maps to a label combination

Disadvantages

- The number of choices in the new class can be exponential ($|Y_{LP}| = O(2^d)$) and Learning a multi-class classifier on exponential choices is expensive

Classifier Chains

- Similar to Binary relevance, doing multiple single-label problem but we use the previous prediction as input feature vector.

X1	X2	Y1
0.3	0.7	1
0.9	0.3	1
0.6	0.6	0
0.4	0.1	0
0.5	0.2	1

Y1 as a class

X1	X2	Y1	Y2
0.3	0.7	1	1
0.9	0.3	1	1
0.6	0.6	0	0
0.4	0.1	0	1
0.5	0.2	1	0

Y1 as a feature and Y2 as a class

X1	X2	Y1	Y2	Y3
0.3	0.7	1	1	0
0.9	0.3	1	1	0
0.6	0.6	0	0	1
0.4	0.1	0	1	0
0.5	0.2	1	0	1

Y1,Y2 as feature and Y3 as a class

Limitation: The result can vary for different order of chains Solution: ensemble

Common Classification Methods

- K-Nearest-Neighbor
- Decision Tree
- Naïve Bayesian
- Support Vector Machine (SVM)
- Artificial Neural Network (ANN)

k-Nearest Neighbour (k-NN) Classification

In *k-nearest-neighbor* (*k*-NN) classification, the training dataset is used to classify each member of a — target dataset.

There is no model created during a learning phase but the training set itself.

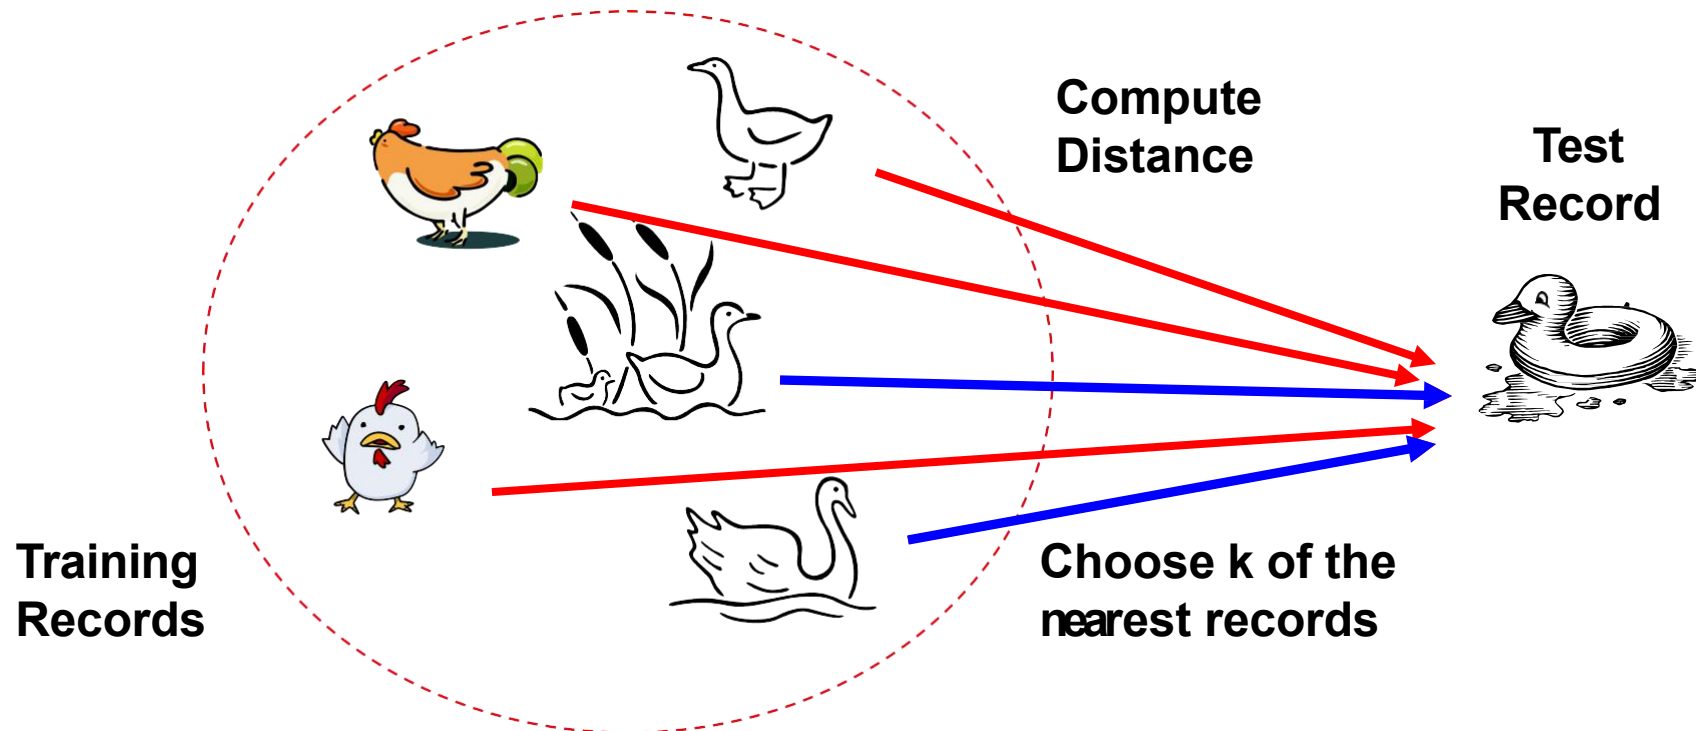
Is known as a *lazy-learning* method unlike *eager-learners* like DT, Naïve Bayes etc.

Rather than building a model and referring to it during the classification, *k*-NN directly refers to the training set for classification.


Nearest Neighbor Classifiers

Basic idea: The basic idea of *nearest-neighbor* models is that the properties of any particular input **X** are likely to be similar to those of points in the neighborhood of **X**


- E.g., If the object walks like a duck, quacks like a duck, then it's probably a duck




Definition of Nearest Neighbor



(a) 1-nearest neighbor



(b) 2-nearest neighbor



(c) 3-nearest neighbor

Three points required to deal with KNN

- The set of stored records

- Distance Metric to compute distance between records

- The value of k , the number of nearest neighbors to retrieve

Nearest Neighbor Classification

- The most used distance function is the Euclidian distance:

$$d(X, Y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

- However, other measures are possible

- the Manhattan distance:

$$d(X, Y) = \sum_{i=1}^n (x_i - y_i)$$

- the Chebychev:

$$d(X, Y) = \max_{i=1}^n (x_i - y_i)$$

- the cosine measure:

$$d(X, Y) = \frac{\sum_{i=1}^n (x_i \cdot y_i)}{\sqrt{\sum_{i=1}^n x_i} \cdot \sqrt{\sum_{i=1}^n y_i}}$$

- Pearson's correlation:


$$d(X, Y) = \frac{\sum_{i=1}^n (x_i - \bar{x}) \cdot (y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \cdot \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

Nearest Neighbor Classification

Choosing the value of k :

If k is too small, sensitive to noise points

If k is too large, neighborhood may include points from other classes



Steps of K -NN Algorithm

k -nearest neighbors algorithm steps.

Step 1. Determine parameter k = number of *nearest neighbors*

Step 2. Calculate the distance between the query-instance and all the training samples

Step 3. Sort the distance and determine *nearest neighbors* based on the k -th minimum distance

Step 4. Gather the category Y of the nearest neighbor

Step 5. Use simple majority of the category of *nearest neighbors* as the prediction value of the query instance

k-NN Example

Question: A factory produces a new paper tissue that passes the laboratory test with $X_1 = 3$ and $X_2 = 7$. Classify this paper as *good* or *bad* using *k*-nearest neighbor method.

Solution

Step 1. Determine the nearest neighbor parameter *k*. In this example we assume *k*=3

Acid Durability (X_1)	Strength (X_2)	Classification (Y)
7	7	Bad
7	4	Bad
3	4	Good
1	4	Good

k-NN Example

Step2. Calculate the distance between the query-instance and all the training examples

Acid Durability (X_1)	Strength (X_2)	Square distance to query instance (3,7)
7	7	$(7-3)^2 + (7-7)^2 = 16$
7	4	$(7-3)^2 + (4-7)^2 = 25$
3	4	$(3-3)^2 + (4-7)^2 = 9$
1	4	$(1-3)^2 + (4-7)^2 = 13$

k-NN Example

Step 3. Sort the distance and determine nearest neighbors based on the *k*-th minimum distance

Acid Durability (X_1)	Strength (X_2)	Square distance to query instance (3,7)	Rank minimum distance	Is it included in 3- nearest neighbors
7	7	$(7-3)^2 + (7-7)^2 = 16$	3	Yes
7	4	$(7-3)^2 + (4-7)^2 = 25$	4	No
3	4	$(3-3)^2 + (4-7)^2 = 9$	1	Yes
1	4	$(1-3)^2 + (4-7)^2 = 13$	2	Yes

k-NN Example

Step 4. Gather the category *Y* of the nearest neighbors.

Notice that the second row last column that the category of nearest neighbor (*Y*) is not included because the rank of this data is more than 3 ($=k$).

Acid Durability (X_1)	Strength (X_2)	Square distance to query instance (3,7)	Rank minimum distance	Is it included in 3-nearest neighbors?	<i>Y</i> = Category of nearest neighbor
7	7	$(7-3)^2 + (7-7)^2 = 16$	3	Yes	Bad
7	4	$(7-3)^2 + (4-7)^2 = 25$	4	No	-
3	4	$(3-3)^2 + (4-7)^2 = 9$	1	Yes	Good
1	4	$(1-3)^2 + (4-7)^2 = 13$	2	Yes	Good

k -NN Example

Step 5. Use simple majority of the category of nearest neighbors as the prediction value of the query instance.

In this example we have 2 good and 1 bad, since $2 > 1$, we conclude that a new paper tissue that pass lab test with $X_1 = 3$ and $X_2 = 7$ is classified as **GOOD**.

* Please try $(x_1=2, x_2=6)$

Nearest Neighbor Classification

- **Advantages of KNN**

- It is extremely easy to implement
- Requires no training prior to making real time predictions. This makes the KNN algorithm much faster than other algorithms that require training e.g SVM, linear regression.
- There are only two parameters required to implement KNN i.e. the value of K and the distance function (e.g. Euclidean or Manhattan etc.)

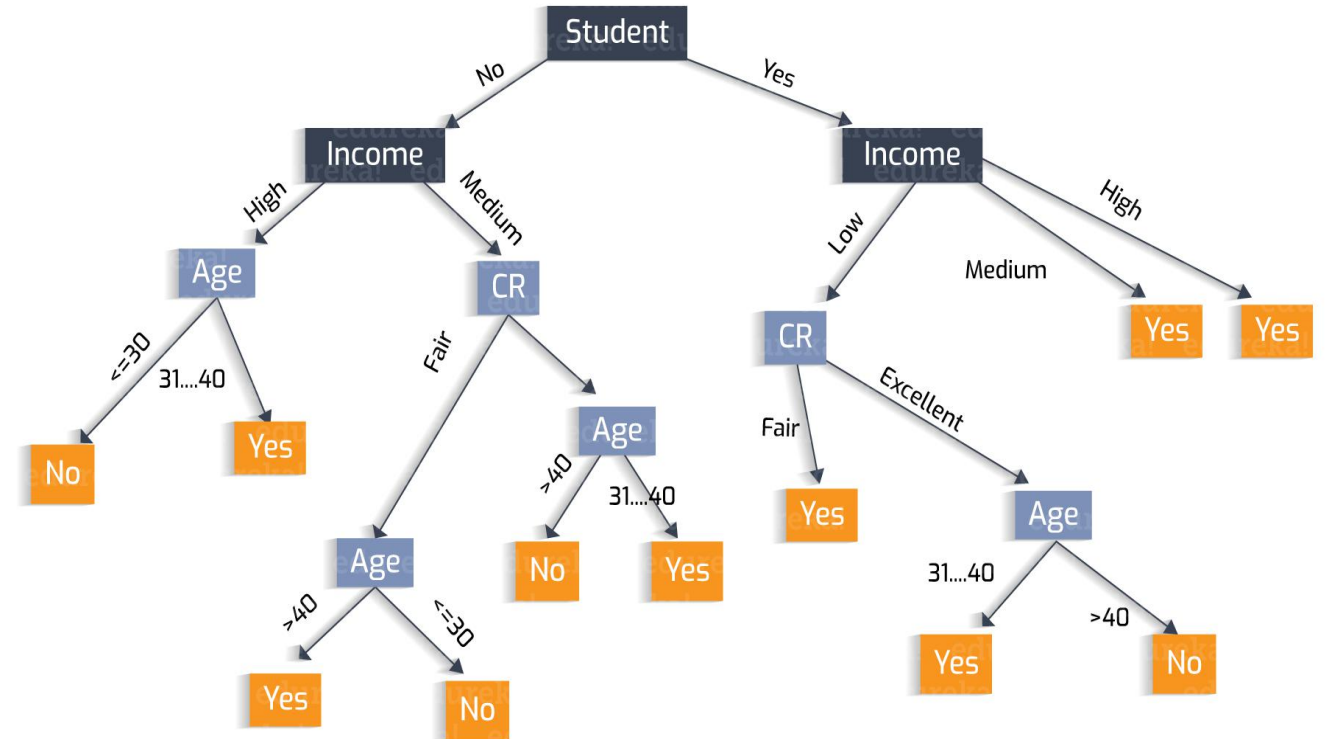

- **Disadvantages of KNN**

- The KNN algorithm doesn't work well with **high dimensional data** because with large number of dimensions, it becomes difficult for the algorithm to **calculate distance in each dimension**.
- The KNN algorithm doesn't work well with **categorical features** since it is difficult to find the distance between dimensions with categorical features.

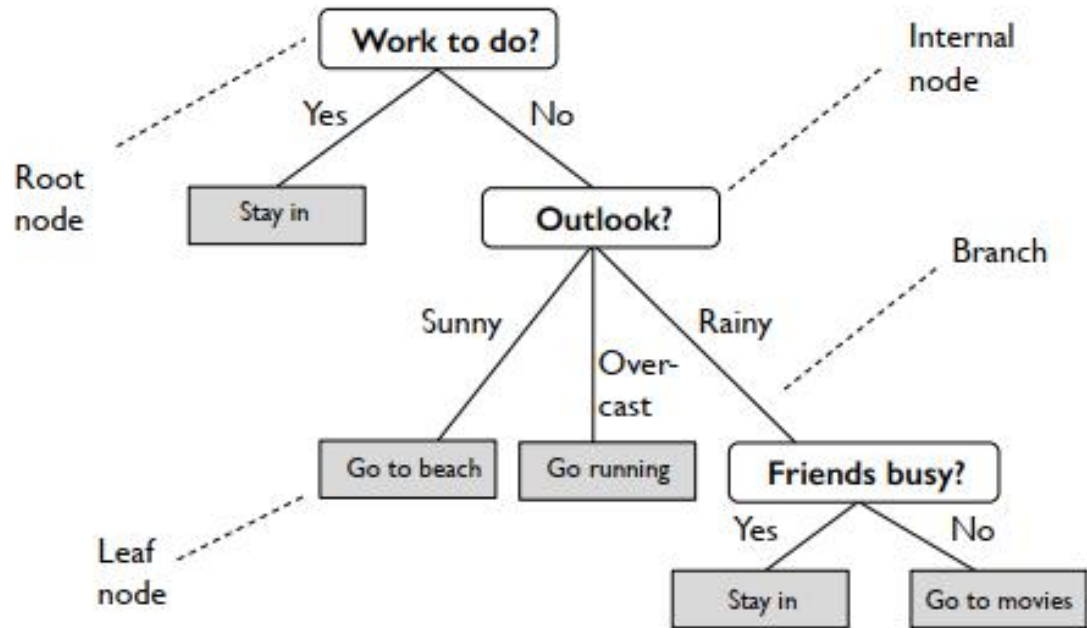
Common Classification Methods

- K-Nearest-Neighbor
- Decision Tree
- Naïve Bayesian
- Support Vector Machine (SVM)
- Artificial Neural Network (ANN)

Decision Tree



Decision Tree



Example of a non-binary decision tree with categorical features.

Decision tree structure

- **Root node:** beginning of a tree and represents **entire population** being analyzed.
- **Internal node:** denotes a test on an **attribute**
- **Branch:** represents an **outcome** of the test
- **Leaf nodes:** represent **class labels** or class distribution

- Tree is constructed in a **top-down recursive divide-and-conquer manner**

Decision Tree

- **Root node:** has **no incoming edges** and zero or more outgoing edges.
- **Internal nodes:** each of which has exactly **one incoming edge** and **two or more** outgoing edges.
- **Leaf** or **terminal node**, each of which has exactly one incoming and no outgoing edges.

Solving the classification problem using DT is a two-step process:

- Decision Tree Induction- Construct a DT using training data/Induction
- For each t_i in D , apply the DT to determine its class/ Deduction


Decision Tree-Example

Output: A Decision Tree for “*buys_computer*”

Training Dataset

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

Decision Tree



Decision Tree

- If there are so many possible trees, can we actually search this space? (solution: greedy search).
 - **Aim**: find a small tree consistent with the training examples
 - **Idea**: (recursively) choose "most significant" attribute as root of (sub)tree.
-
- A **greedy algorithm** is a simple, intuitive **algorithm** that is used in optimization problems. The **algorithm** makes the optimal choice at each step as it attempts to find the overall optimal way to solve the entire problem.

Decision Tree

Algorithm:

- The algorithm is a greedy algorithm that doesn't warranty optimality
- The algorithm for decision tree involves three general phases as stated below:
- **Phase I:** Find Splitting Criteria based on all the sample set at the splitting point (node) (*attribute selection measure*)
- **Phase II:** Split all the sample data based on the splitting criteria and form branches and each successor nodes will have the samples
- **Phase III:** Do phase one and phase two iteratively until stopping criterion get fulfilled

Decision Tree

Attribute Selection Measure

1. Information Gain (ID3 algorithm)
 2. Gain Ratio (C4.5 algorithm) [will not be discussed]
 3. Gini Index (CART algorithm) [will not be discussed]
- ID3, C4.5, and CART adopt a greedy (i.e., nonbacktracking) approach in which decision trees are constructed in a top-down recursive divide-and-conquer manner using Information gain attribute selection approach.

Decision Tree

Information Gain (ID3)

- Developed by J. Ross Quinlan in late 1970s
- All attributes are assumed to be categorical
- Uses information gain to split node into branches
- Details of the algorithm is given below
 - Select the **attribute** with the highest **information gain**
 - This attribute minimizes the information needed to classify the tuples in the resulting partitions and reflects the least randomness or “impurity” in these partitions
 - Such an approach minimizes the expected number of tests needed to classify a given tuple and guarantees that a simple (but not necessarily the simplest) tree is found.
 - The amount of information need becomes maximum if all the classes has the same number of tuples (i.e. $I(D)=1$)
 - The amount of information need becomes minimal if all the tuples belongs to one class (i.e $I(D)= 0$)

Decision Tree

Information Gain (ID3)

- Let p_i be the probability that an arbitrary tuple in D belongs to class C_i , estimated by $|C_{i,D}|/|D|$, and $|D_j|$ is those tuples in D that have outcome a_j of A for $i=1$ to m , and $j=1$ to v .

Expected information (entropy) needed to classify a tuple in D :

$$Info(D) = -\sum_{i=1}^m p_i \log_2(p_i)$$

- $info(D)$ is average amount of information needed to identify the class label of a tuple in D . $Info(D)$ is also known as **entropy** of D .
- Information after split (after using A to split D into v partitions) to classify D :

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times I(D_j)$$

- Information gained by branching on attribute A

$$Gain(A) = Info(D) - Info_A(D)$$

Decision Tree

Information Gain (ID3)-Example

■ Class P: buys_computer = "yes"

■ Class N: buys_computer = "no"

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940$$

age	p _i	n _i	I(p _i , n _i)
<=30	2	3	0.971
31...40	4	0	0
>40	3	2	0.971

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

$\frac{5}{14} I(2,3)$ means "age <=30" has 5 out of 14 samples, with 2 yes'es and 3 no's. Hence

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

Similarly,

$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

$$Gain(credit_rating) = 0.048$$

Decision Tree

Stopping Criteria

- Some conditions for stopping partitioning
 - All samples for a given node belong to the same class
 - There are no remaining attributes for further partitioning
 - There are no samples left in the branching direction
 - Some user defined stopping criterion met such as
 - Depth threshold
 - Number of samples in the node becomes below some value
 - Or others

Decision Tree

Classification Rule Extraction

- Represent the knowledge in the form of **IF-THEN** rules
- One rule is created for each path from the root to a leaf
- Each attribute-value pair along a path forms a conjunction
- The leaf node holds the class prediction
- Rules are easier for humans to understand
- Example
 - IF *age* = “≤30” AND *student* = “no” THEN *buys_computer* = “no”
 - IF *age* = “≤30” AND *student* = “yes” THEN *buys_computer* = “yes”
 - IF *age* = “31...40” THEN *buys_computer* = “yes”
 - IF *age* = “>40” AND *credit_rating* = “excellent” THEN *buys_computer* = “yes”
 - IF *age* = “>40” AND *credit_rating* = “fair” THEN *buys_computer* = “no”

Decision Tree

Strength

- In practice: One of the most popular methods. Why?
 - Very comprehensible – the tree structure specifies the entire decision structure
 - Easy for decision makers to understand model's rational
 - Relatively easy to implement
 - Very fast to run (to classify examples) with large data sets


Weakness

- Splits are simple and not necessarily optimal
- Overfitting (may not have the same performance when deployed, complex and more specific conditions)
- Underfitting (may over generalize during training, less conditions)

Decision Tree

Weaknesses: Overfitting


Overfitting means that the model performs poorly on new examples (e.g. *testing* examples) as it is too highly trained to the specific (non-general) nuances of the *training* examples.



Decision Tree

Weaknesses: Overfitting


- Learning a tree that classifies the training data perfectly may not lead to the tree with the best generalization to unseen data.
 - There may be noise in the training data that the tree is erroneously fitting.
 - The algorithm may be making poor decisions towards the leaves of the tree that are based on very little data and may not reflect reliable trends.
- A hypothesis, h , is said to overfit the training data if there exists another hypothesis h' such that h has less error than h' on the training data but greater error on independent test data.



Decision Tree


Weaknesses: Underfitting

Underfitting means that the model performs poorly on new examples as it is too simplistic to distinguish between them (i.e. has not picked up the important patterns from the training examples)



Bayesian Classification

$$P(h | X) = \frac{P(X | h) P(h)}{P(X)}$$



Thomas Bayes 1702-1761

Bayesian Classification

- Bayesian classifiers are statistical classifiers.
- They can predict class membership probabilities, such as **the probability** that a **given tuple belongs** to a particular class.
- Bayesian classification is based on **Bayes' theorem**

$$P(h | X) = \frac{P(X | h)P(h)}{P(X)}$$

where h is hypothesis, X is a training data

- Compared to Decision tree, Bayesian classifiers have also exhibited high accuracy and speed when applied to large databases.

Bayesian Classification

Baye's Theroem

- Let **X** be an observation (vector of values of the independent variables) whose class label (dependent variable) is not given
- Let **h** be a hypothesis that says the class of **X** is **C**
- The classification problem require to accept or reject the hypothesis
- This can be done in Bayesian classification by finding the *posteriori probability of a hypothesis* $P(h | X)$
- Bayes theorem says
$$P(h | X) = \frac{P(X | h)P(h)}{P(X)}$$

Naïve Bayesian Classification

- 1) Let D be a training set of tuples and their associated class labels, and each tuple is represented by an n -dimensional feature vector $X = (x_1, x_2, \dots, x_n)$, depicting n measurements made on the tuple from n attributes, A_1, A_2, \dots, A_n , respectively.
- 2) And suppose there are m classes C_1, C_2, \dots, C_m . Now, given a tuple X , the classifier predict that X belongs to the class having highest posterior probability (conditioned on X .)
- 3) That is, classifier assigns an unknown sample X to the class C_i , if and only if $P(C_i|X) > P(C_j|X)$ for $1 \leq j \leq m, j \neq i$. The class C_i for which $P(C_i|X)$ is maximized is called the **maximum posteriori hypothesis**.

This can be derived from Bayes' theorem:

$$P(C_i | X) = \frac{P(X|C_i)P(C_i)}{P(X)}$$

- 4) Since $P(X)$ is constant for all classes, only $P(C_i|X) = P(X|C_i)P(C_i)$ needs to be maximized. If prior probability of X unknown, classes are assumed to be equally likely, i.e., $P(C_1) = P(C_2) = P(C_3) = \dots = P(C_m)$ and would maximize $P(X|C_i)$, where $P(C_i) = |C_{i,D}| / |D|$.

Naïve Bayesian Classification

- 5) A simplified assumption: attributes are conditionally independent (i.e., no dependence relation between attributes):

$$P(\mathbf{X}|C_i) = \prod_{k=1}^n P(x_k | C_i) = P(x_1 | C_i) \times P(x_2 | C_i) \times \dots \times P(x_n | C_i)$$

- 6) This greatly reduces the computation cost:-only counts the class distribution
- 7) If A_k is categorical, $P(x_k|C_i) = s_{ik}/s_i$, is the # of tuples of C_i in D having value x_k for A_k divided by $|C_i, D|$ (# of tuples of C_i in D)

S

- 8) If A_k is continuous-valued, $P(x_k|C_i)$ is usually computed based on Gaussian distribution with a mean μ and standard deviation σ

$$g(x, \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

and $P(x_k|C_i)$ is $P(\mathbf{X}|C_i) = g(x_k, \mu_{C_i}, \sigma_{C_i})$

- 9) In order to classify an unknown sample X , $P(X|C_j)P(C_j)$ is evaluated for each class C_i . X is assigned to the class C_i iff, $P(X|C_i)P(C_i) > P(X|C_j)P(C_j)$ for $1 \leq j \leq m, j \neq i$.

Naïve Bayesian Classification

Example:

Class:

C1:buys_computer = "yes"

C2:buys_computer = "no"

Given a data sample X:

X = (age <=30, Income = medium, Student = yes, Credit_rating = Fair)

Task:

Classify **X** using Bayesian classifier

age	income	student	redit_rating	_com
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

Naïve Bayesian Classification

Example:

We need to maximize $P(X|C_i)P(C_i)$, for $i=1,2$.

$P(C_i)$: $P(\text{buys_computer} = \text{"yes"}) = 9/14 = 0.643$

$P(\text{buys_computer} = \text{"no"}) = 5/14 = 0.357$

Compute $P(X|C_i)$ for each class:

$P(\text{age} = \text{"<=30"} | \text{buys_computer} = \text{"yes"}) = 2/9 = 0.222$

$P(\text{age} = \text{"<= 30"} | \text{buys_computer} = \text{"no"}) = 3/5 = 0.6$

$P(\text{income} = \text{"medium"} | \text{buys_computer} = \text{"yes"}) = 4/9 = 0.444$

$P(\text{income} = \text{"medium"} | \text{buys_computer} = \text{"no"}) = 2/5 = 0.4$

$P(\text{student} = \text{"yes"} | \text{buys_computer} = \text{"yes"}) = 6/9 = 0.667$

$P(\text{student} = \text{"yes"} | \text{buys_computer} = \text{"no"}) = 1/5 = 0.2$

$P(\text{credit_rating} = \text{"fair"} | \text{buys_computer} = \text{"yes"}) = 6/9 = 0.667$

$P(\text{credit_rating} = \text{"fair"} | \text{buys_computer} = \text{"no"}) = 2/5 = 0.4$

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

Naïve Bayesian Classification

Example:

$X = (\text{age} \leq 30, \text{income} = \text{medium}, \text{student} = \text{yes}, \text{credit_rating} = \text{fair})$

- $P(X|C_i) : P(X|\text{buys_computer} = \text{"yes"})$
 $= 0.222 \times 0.444 \times 0.667 \times 0.667$
 $= 0.044$
- $P(X|\text{buys_computer} = \text{"no"})$
 $= 0.6 \times 0.4 \times 0.2 \times 0.4$
 $= 0.019$
- $P(X|C_i) * P(C_i) : P(X|\text{buys_computer} = \text{"yes"}) * P(\text{buys_computer} = \text{"yes"})$
 $= 0.044 * 0.643$
 $= 0.028$
- $P(X|\text{buys_computer} = \text{"no"}) * P(\text{buys_computer} = \text{"no"})$
 $= 0.019 * 0.357$
 $= 0.007$
- Since $0.028 > 0.007$, **X** belongs to class (**"buys_computer = yes"**)

age	income	student	redit_rating	_com
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

Naïve Bayesian Classification

Naïve Bayesian prediction requires each conditional prob. be **non-zero**. Otherwise, the predicted prob. will be zero

$$P(X \mid C_i) = \prod_{k=1}^n P(x_k \mid C_i)$$

Ex. Suppose a dataset with 1000 tuples, income=low (0), income= medium (990), and income = high (10)

Use **Laplacian correction** (or Laplacian estimator)

- *Adding 1 to each case*

Prob(income = low) = 1/1003

Prob(income = medium) = 991/1003

Prob(income = high) = 11/1003

Naïve Bayesian Classification

Naïve Bayes Classifier: Strength


- Simple
- Incremental learning
- Naturally a probability estimator
- Easily handles missing values

Naïve Bayesian Classification

Naïve Bayes Classifier: Weakness

- Independence assumption
- Categorical/discrete attributes
- Sensitive to missing values


Support Vector Machine



Support Vector Machine




- SVM is a geometric model that views the input data as **two sets of vectors** in an n -dimensional space.
- It constructs a **separating hyperplane** in that space, one which maximizes the *margin* between the two data sets.
- To calculate the margin, **two parallel hyperplanes** are constructed, one on each side of the separating hyperplane.
- A good separation is achieved by the hyperplane that has the **largest distance** to the neighboring data points of both classes.
- The vectors (points) that constrain the width of the margin are the ***support vectors***.

Support Vector Machine




An SVM analysis finds the line (or, in general, hyperplane) that is oriented so that the margin between the support vectors is maximized. In the figure above, Solution 2 is superior to Solution 1 because it has a larger margin.

Support Vector Machine





- To separate the two classes of data points:
 - There are many possible hyperplanes that could be chosen.
 - Our objective is to find a plane that has the maximum margin, i.e the maximum distance between data points of both classes.
 - Maximizing the margin distance provides some reinforcement so that future data points can be classified with more confidence.

Support Vector Machine



SVM algorithm has a feature to ignore **outliers** and find the hyper-plane that has the maximum margin




Add a new feature $z=x^2+y^2$ and then plot the data points on axis x and z

kernel trick. The SVM kernel is a function that takes low dimensional input space and transforms it to a higher dimensional space i.e. it converts not separable problem to separable problem

Support Vector Machine – Kernel Functions

- Kernel function Φ : map data into a different space to enable linear separation.



- Convert linear learning machines (such as linear SVM) into non-linear ones (such as kernel SVM) by doing an inner product between datapoints.
- Kernel functions are very powerful. They allow SVM models to perform separations even with very complex boundaries.
 - Some popular kernel functions are linear, polynomial, and radial basis.
 - While you can construct your own kernel functions according to the data structure, ML tool provides a variety of built-in kernels. e.g [Scikit-learn](#)

Support Vector Machine – Multi-class classification

- Suppose we have a bunch of source code which are going to be classified in 3 programming languages. Do you think the binary classifier can solve this problem?

```
for (int i = 0; i < 5; i++)  
{  
    cout << i << "\n";  
}
```

1

```
for x in range(0,5):  
    print(x)
```


2

```
for (int i = 0; i < 5; i++)  
{  
    cout << i ;  
}
```

3

```
for (int i = 0; i < 5; i++)  
{  
    System.out.println(i);  
}
```

4



- One-Vs-Rest for Multi-Class Classification: distinguishing between some label and all the others, where the class prediction with highest probability wins.
- One-Vs-One for Multi-Class Classification: a classifier is trained for every pair of classes, making continuous comparisons. The class prediction with highest quantity of predictions wins.

Support Vector Machine – Multi-class classification

- Suppose we have a bunch of source code which are going to be classified in 3 programming languages. Do you think the binary classifier can solve this problem?

```
for (int i = 0; i < 5; i++)  
{  
    cout << i << "\n";  
}
```

1

```
for x in range(0,5):  
  
    print(x)
```

2

```
for (int i = 0; i < 5; i++)  
{  
    cout << i ;  
}
```

3

```
for (int i = 0; i < 5; i++)  
{  
    System.out.println(i);  
}
```

4



- In case of OvR: creating C binary classifiers for C number of class, where each binary classifier should predict a class probability as p_1, \dots and p_C . Then the prediction will be $\text{argmax}(p_1, p_2, p_3)$
 - binary classifier 1: java vs {c++, python}
 - binary classifier 2: c++ vs {java, python}
 - binary classifier 3: python vs {c++, java}
- if the probability that **source code #4** is Java vs c++ or python is 0.6, c++ vs java or python is 0.15, python vs c++ or java is 0.25 then then $\text{argmax}(0.6, 0.15, 0.25)=0$ which is JAVA

Support Vector Machine – Multi-class classification

- Suppose we have a bunch of source code which are going to be classified in 3 programming languages. Do you think the binary classifier can solve this problem?

```
for (int i = 0; i < 5; i++)  
{  
    cout << i << "\n";  
}
```

1

```
for x in range(0,5):  
    print(x)
```

2

```
for (int i = 0; i < 5; i++)  
{  
    cout << i ;  
}
```

3

```
for (int i = 0; i < 5; i++)  
{  
    System.out.println(i);  
}
```

4



- In case of OvO: creating set of binary classifiers in $(C(C-1)/2)$ for C number of class, each representing one of the pairs. Then the class that is picked the most is the winner
 - binary classifier 1: java vs c++
 - binary classifier 2: java vs python
 - binary classifier 3: python vs c++
- if **source code #4** is classified as Java in the first and the second case, it is automatically JAVA

- Practical Examples

Support Vector Machine – Example

```
from sklearn import svm
import numpy as np
x_train=np.array([[4],[5],[6],[7],[8]])
y_train=np.array([4,5,6,7,8])
est = svm.LinearSVC()

#training
Model=est.fit(x_train, y_train)


#prediction
pred = est.decision_function([[-5], [6], [7],[8]])
```

* If none of the predictions get more than half of the **votes**, we may say that the ensemble method **could not make a stable prediction for this instance**.

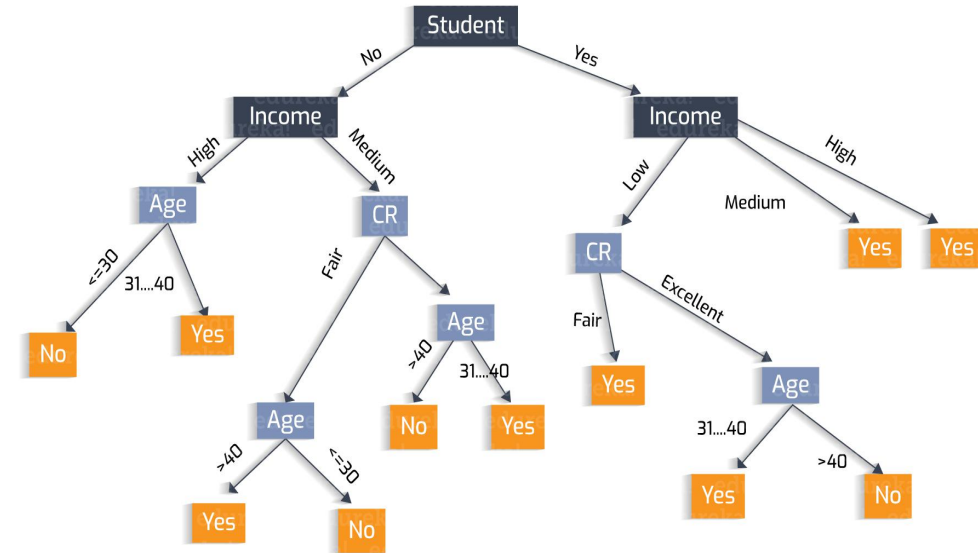
Classification Algorithm: Others

- Please review the following methods
 - **Classification**
 - Bayesian Belief Networks
 - case-based reasoning
 - Genetic algorithm and programming
 - Fuzzy set approaches
 - **Regression**
 - Linear regression
 - Ridge regression

age	income	student	credit_rating	con
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no



$$P(h|X) = \frac{P(X|h)P(h)}{P(X)}$$



```
for (int i = 0; i < 5; i++)
{
    cout << i << "\n";
}
```

1

```
for x in range(0,5):
    print(x)
```

2

```
for (int i = 0; i < 5; i++)
{
    cout << i ;
}
```

3

```
for (int i = 0; i < 5; i++)
{
    System.out.println(i);
}
```

4

