




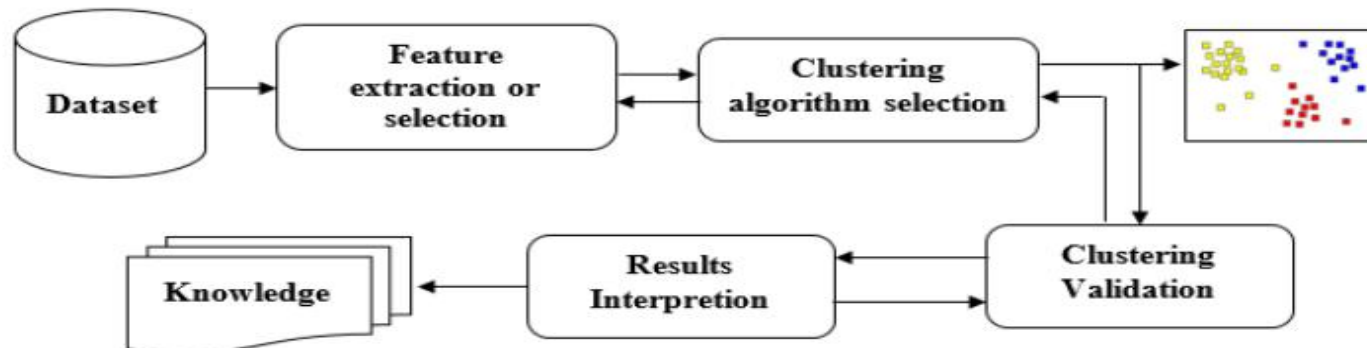

Chapter-6

Unsupervised Learning

Unsupervised Learning in ML




Dimensionality Reduction & Principal Component Analysis



Unsupervised learning

- Unsupervised learning aims to find the underlying structure or the distribution of data. We want to explore the data to find some intrinsic structures in them.



- Learn from inputs $x_1, \dots, x_n \in \mathbb{R}^d$ without any labels y_1, \dots, y_n .

***No predefined classes - Clustering**

What is Cluster Analysis?


- **Cluster**: a collection of data objects
 - Similar to one another within the same cluster
 - Dissimilar to the objects in other clusters
- **Cluster analysis**

Finding similarities between data according to the characteristics found in the data and grouping similar data objects into clusters

Definition. Given a database $D = \{t_1, t_2, \dots, t_n\}$ of tuples and an integer value k , the **clustering problem** is to define mapping $f: D \rightarrow \{1, \dots, k\}$ where each t_i is assigned to one cluster K_j , $1 \leq j \leq k$. A **cluster**, K_j , containing precisely those tuples mapped to it; that is, $K_j = \{t_i \mid f(t_i) = K_j, 1 \leq i \leq n, \text{ and } t_i \in D\}$

What is Cluster Analysis?

Finding groups of objects such that the objects in a group will be similar (or related) to one another and different from (or unrelated to) the objects in other groups



Clustering-Applications

- **Recommender systems**: organizing products and customers into groups that are similar



- **Social networks**: cluster users into groups that have similar interests/preferences

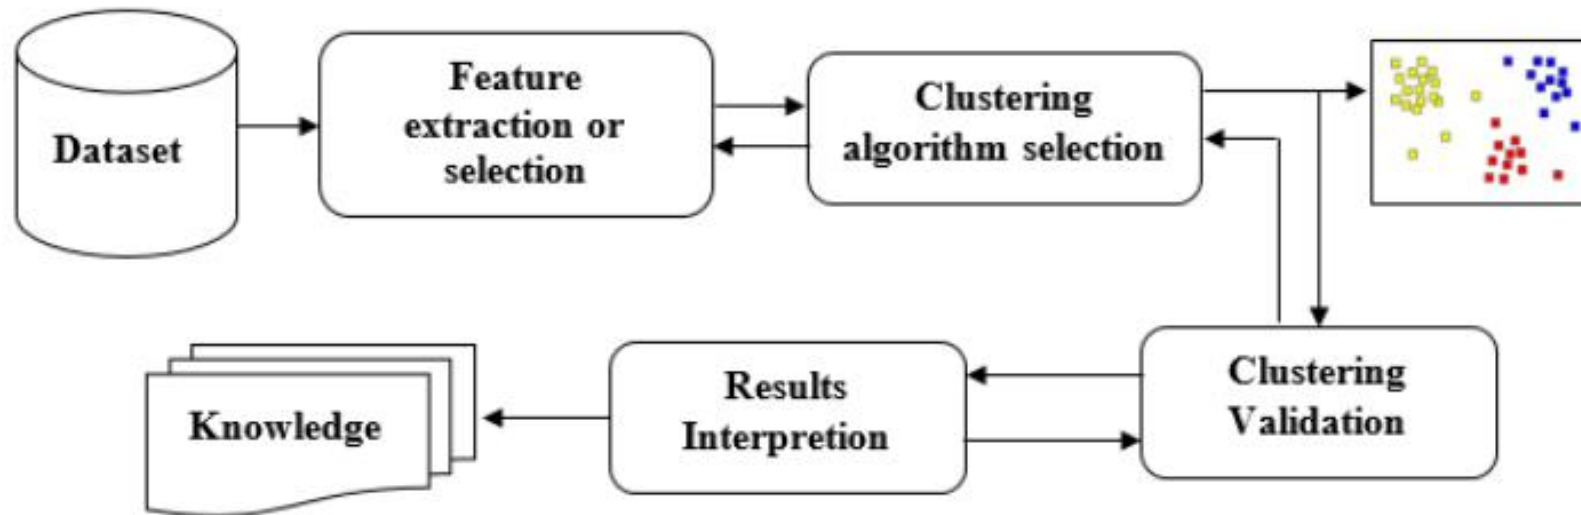


- **WWW document classification**: organize webpages (e.g., news articles) into clusters with similar content (sports, economics, ...)



Procedures of Cluster Analysis

- A typical cluster analysis consists of four steps



Procedures of Cluster Analysis

Feature extraction

- Feature extraction is the process of using one or more transformations of the input features to generate new principal features.
- Feature extraction can be elaborated in the context of dimensionality reduction and data visualization.

Procedures of Cluster Analysis

Clustering algorithm design or selection

- The clustering step is usually combined with the selection of a corresponding proximity (i.e, the closeness or distance) measure and
- the construction of a clustering criterion function (i.e, finding the optimal partitioning of a data set according to some criterion function or algorithm).
- **Proximity measures** refer to the measures of Similarity and Dissimilarity.
- **Clustering criterion:** once a proximity measure is chosen, the construction of a clustering criterion function makes the partition of clusters as an optimization problem.

Procedures of Cluster Analysis

Cluster validation (Assessment of results)

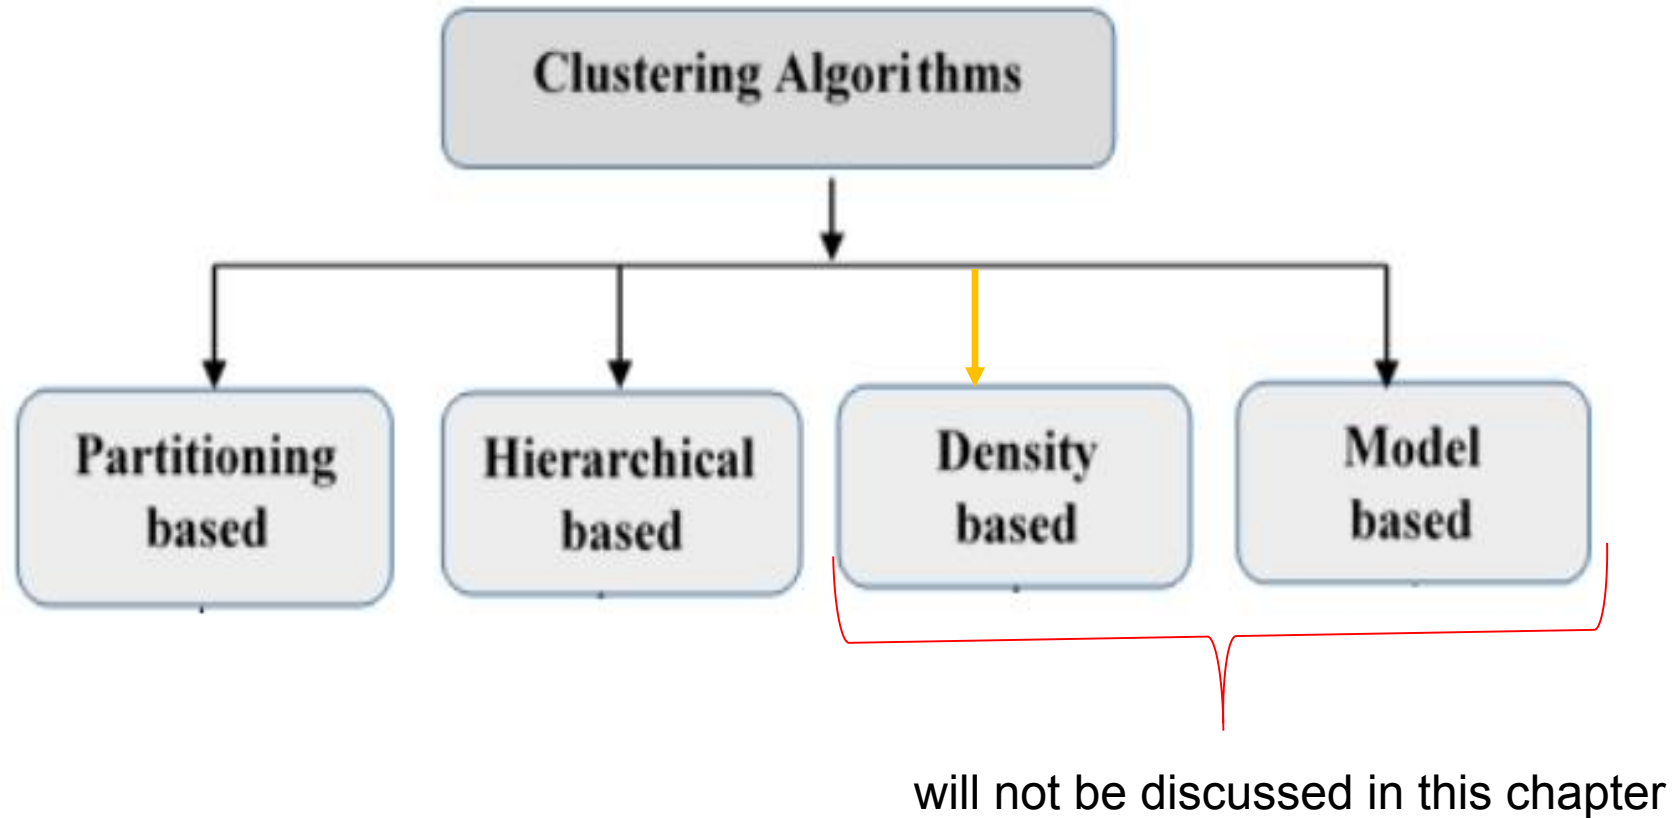
- Given a dataset, any clustering algorithm can usually generate clusters, no matter whether the structure exists or not.
- Thus, effective evaluation standards and criteria are essential to provide the users with a degree of confidence for the clustering results.
- These evaluation methods should be objective and have no preference for any algorithm.

Procedures of Cluster Analysis

Results Interpretation

- The final target of clustering is to supply users with meaningful perceptions from the original dataset, with the aim that they can effectively solve the problems faced.
- Experts in different domains interpret the data groupings.
- Further analyses, even experiments, may be required to assure the reliability of extracted knowledge.

Categories of clustering Algorithms



Partitioning method

Partitioning method: Construct a partition of a database **D** of **n** objects into a set of **k** clusters

Various approaches have been proposed and some of them are:

K-mean Approach

K-medoid Approach, CLARA (Custering LARge Application) and CLARANS (CLustering Algorithm based on RANdomized Search) [**will not be covered here**]

K-means clustering

- K-means is a **partitional clustering** algorithm
- Let the set of data points (or instances) D be

$$\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\},$$

where $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{ir})$ is a **vector** in a real-valued space $X \in R^r$, and r is the number of attributes (dimensions) in the data.

- The k -means algorithm partitions the given data into k clusters.
 - Each cluster has a cluster **center**, called **centroid**.
 - k is specified by the user

K-means algorithm

Given k , the *k-means* algorithm is implemented in 4 steps:

- Select K centroid (Can be K values randomly, or K data points randomly)
- Partition objects into k subsets. An object will be clustered into class J if it has the smallest distance with this class mean compared to the distance with the other class mean
- Compute the new centroids of the clusters of the current partition. The centroid of the j th cluster is the center (mean point) of the data point whose cluster index is found to be the center of class j in the above step.
- Go back to Step 3, stop when the process converge.

K-means algorithm

```
1: function KMEANS(parameter  $k$ , inputs  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ )
2:   initialize cluster centers  $\mathbf{c}_1, \dots, \mathbf{c}_k$ 
3:   repeat
4:     for  $i = 1 : n$  do
5:       label the input  $\mathbf{x}_i$  as belonging to the nearest cluster,
          
$$y_i := \arg \min_{j=1, \dots, k} \|\mathbf{x}_i - \mathbf{c}_j\|^2$$

6:     end for
7:     for  $j = 1 : k$  do
8:       compute cluster center  $\mathbf{c}_j$  as the mean of all inputs of the  $j$ th cluster,
          
$$\mathbf{c}_j := \text{mean}(\{\mathbf{x}_i : y_i = j\})$$

9:     end for
10:    until convergence criterion is met
11:    return cluster centers  $\mathbf{c}_1, \dots, \mathbf{c}_k$ 
12: end function
```


K-Means Clustering – Simple Example

Example: Suppose we are given the following items to cluster: $\{2, 4, 10, 12, 3, 20, 30, 11, 25\}$ and suppose $k=2$.

Initially assign the means to the first two values: $m_1=2$ and $m_2=4$.

Using Euclidean distance, find initial clusters $K_1 = \{2, 3\}$ and $K_2 = \{4, 10, 12, 20, 30, 11, 25\}$

- For instance, $\text{dist}(2, 10) = \sqrt{(10-2)^2} = 8$; $\text{dist}(4, 10) = \sqrt{(10-4)^2} = 6$, hence item 10 will be assigned to cluster K_2
- But, note that, the item 3 is equally close to both means, so we arbitrarily choose K_1 .

Now compute new means $m_1 = 2.5$ and $m_2 = 16$

Assign items closer to m_1 and m_2 , to get clusters $K_1 = \{2, 3, 4\}$ and $K_2 = \{10, 12, 20, 30, 11, 25\}$

K-Means Clustering – Simple Example

- Continue in this fashion till the values of K_1 and K_2 do not change as shown in the table.
- Since there is no change in K_1 and K_2 (last two rows of the table), the answer is : $K_1 = \{2, 3, 4, 10, 11, 12\}$ and $K_2 = \{20, 30, 25\}$

m_1	m_2	K_1	K_2
3	18	$\{2,3,4,10\}$	$\{12,20,30,11,25\}$
4.75	19.6	$\{2,3,4,10,11,12\}$	$20,30,25\}$
7	25	$\{2,3,4,10,11,12\}$	$20,30,25\}$

* No (or minimum) re-assignments of data points to different clusters, or no change of centroids

Strengths of k-means

- Strengths:
 - Simple: easy to understand and to implement
 - Efficient: Time complexity: $O(tkn)$, where n is the number of data points, k is the number of clusters, and t is the number of iterations.
 - Since both k and t are small. k -means is considered a linear algorithm.
- K-means is the most popular clustering algorithm.

Weaknesses of k-means



- The algorithm is only applicable if the **mean** is defined.
 - For categorical data, **k-mode** - the centroid is represented by most frequent values.
- The user needs to specify **k**.
- The algorithm is sensitive to **outliers**
 - Outliers are data points that are very far away from other data points.
 - Outliers could be errors in the data recording or some special data points with very different values.

K-means summary

- Despite weaknesses, *k*-means is still the most popular algorithm due to its simplicity, efficiency and
 - other clustering algorithms have their own lists of weaknesses.
- No clear evidence that any other clustering algorithm performs better in general
 - although they may be more suitable for some specific types of data or applications.
- Comparing different clustering algorithms is a difficult task. No one knows the correct clusters!

Hierarchical Clustering

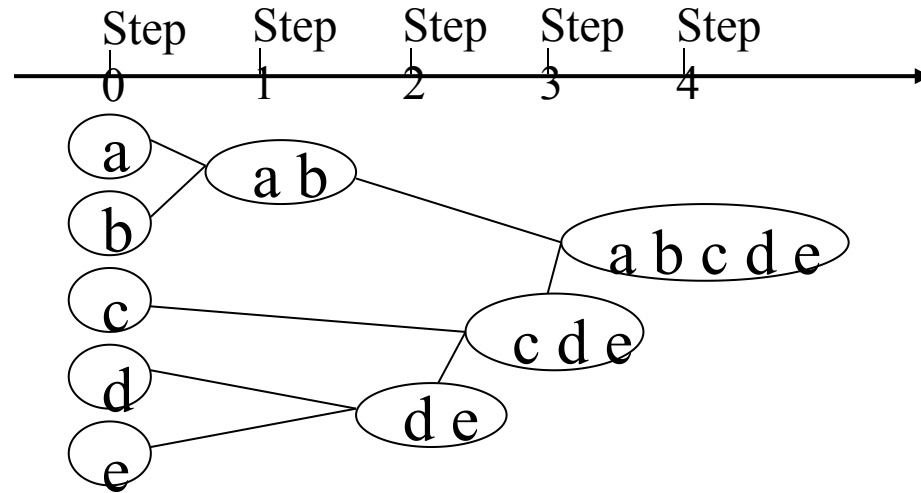
- Produce a nested sequence of clusters, a **tree**, also called **Dendrogram**.



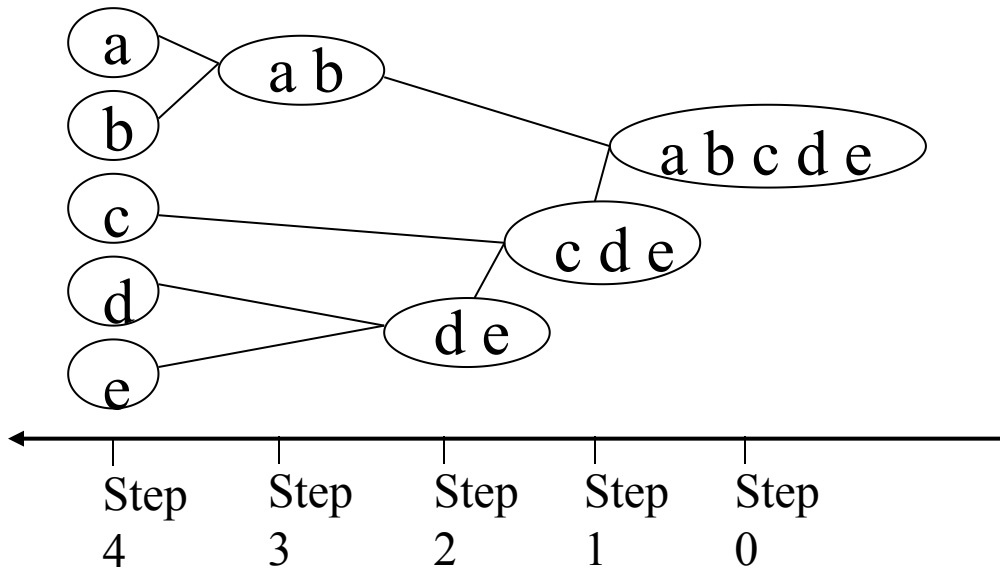
Types of hierarchical clustering

- **Agglomerative (bottom up) clustering**: It builds the dendrogram (tree) from the bottom level, and
 - merges the most similar (or nearest) pair of clusters
 - stops when all the data points are merged into a single cluster (i.e., the root cluster).
- **Divisive (top down) clustering**: It starts with all data points in one cluster, the root.
 - Splits the root into a set of child clusters. Each child cluster is recursively divided further
 - stops when only singleton clusters of individual data points remain, i.e., each cluster with only a single point

Types of hierarchical clustering



**Agglomerative Nesting
(AGNES)**



**Divisive Analysis
(DIANA)**

Agglomerative clustering

It is more popular than divisive methods.



- At the beginning, each data point forms a cluster (also called a node).
- Merge nodes/clusters that have the least distance.
- Go on merging
- Eventually all nodes belong to one cluster

Agglomerative clustering algorithm

```


1: function HIERARCHICALCLUSTERING(inputs  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ )
2:   assign each input to a cluster
3:   repeat
4:     link the two clusters with minimal distance
5:   until only a single root cluster left
6:   return tree of cluster linkages
7: end function

```




Measuring the distance of two clusters


- Simple linkage



- Average linkage




- Complete linkage



- ▶ Let $S_j \subseteq \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ be the set of inputs contained in the j th cluster
- ▶ **Simple linkage:** $d(i, j) := \min_{\mathbf{x} \in S_i, \tilde{\mathbf{x}} \in S_j} \|\mathbf{x} - \tilde{\mathbf{x}}\|$
- ▶ **Average linkage:** $d(i, j) := \text{mean}_{\mathbf{x} \in S_i, \tilde{\mathbf{x}} \in S_j} \|\mathbf{x} - \tilde{\mathbf{x}}\|$
- ▶ **Complete linkage:** $d(i, j) := \max_{\mathbf{x} \in S_i, \tilde{\mathbf{x}} \in S_j} \|\mathbf{x} - \tilde{\mathbf{x}}\|$

Single link method


- The distance between two clusters is the distance between two **closest data points** in the two clusters, one data point from each cluster.
- It can find arbitrarily shaped clusters, but
 - It may cause the undesirable “**chain effect**” by noisy points



Distance between clusters = $d_{2,3}$

Average link and centroid methods

- **Average link**: the distance between two clusters is the average distance of all pair-wise distances between the data points in two clusters.




$$\text{Distance between clusters} = \frac{d_{1,3} + d_{1,4} + d_{1,5} + d_{2,3} + d_{2,4} + d_{2,5}}{6}$$

- **Centroid method**: In this method, the distance between two clusters is the distance between their centroids

Complete link method

- The distance between two clusters is the distance of two **furthest** data points in the two clusters.
- It is sensitive to outliers because they are far away



Distance between clusters = $d_{1,5}$

Hierarchical Clustering: the complexity

- All the algorithms are at least $O(n^2)$. n is the number of data points.
- Single link can be done in $O(n^2)$.
- Complete and average links can be done in $O(n^2 \log n)$.
- Due the complexity, hard to use for large data sets.

Hierarchical Clustering: Problems and Limitations

Once a decision is made to combine two clusters, it cannot be undone

No objective function is directly minimized

Different schemes have problems with one or more of the following:

- Sensitivity to noise and outliers

- Difficulty handling different sized clusters and convex shapes

- Breaking large clusters

How to choose a clustering algorithm

- Clustering research has a long history. A vast collection of algorithms are available.
 - We only introduced several main algorithms.
- **Choosing the “best” algorithm is a challenge.**
 - Every algorithm has limitations and works well with certain data distributions.
 - It is very hard, if not impossible, to know what distribution the application data follow. The data may not fully follow any “ideal” structure or distribution required by the algorithms.
 - One also needs to decide how to standardize the data, to choose a suitable distance function and to select other parameter values.

Choose a clustering algorithm (cont ...)

- Due to these complexities, the common practice is to
 - run several algorithms using different distance functions and parameter settings, and
 - then carefully analyze and compare the results.
- The interpretation of the results must be based on insight into the meaning of the original data together with knowledge of the algorithms used.
- Clustering is highly **application dependent** and to certain extent **subjective** (personal preferences).

Cluster Evaluation

- For cluster analysis, the question is how to evaluate the “goodness” of the resulting clusters?
- But “clusters are in the eye of the beholder”!
- Then why do we want to evaluate them?
 - To avoid finding patterns in noise
 - To compare clustering algorithms
 - To compare two sets of clusters
 - To compare two clusters
 - Measuring Cluster Quality:
 - Ground truth is the ideal clustering that is often built using human experts.
 - **Extrinsic Method:** If ground truth is available, compares the clustering against the ground truth and measure.
 - **Intrinsic Method:** If ground truth is not available, evaluate goodness of a clustering by considering how well the clusters are separated.

Summary

- The validation of clustering structures is the most difficult and frustrating part of cluster analysis.
- Clustering is has along history and still active
 - There are a huge number of clustering algorithms
 - More are still coming every year.
- We only introduced several main algorithms. There are many others, e.g.,
 - sub-space clustering, scale-up methods, neural networks based methods, fuzzy clustering, co-clustering, etc.
- Clustering is hard to evaluate, but very useful in practice. This partially explains why there are still a large number of clustering algorithms being devised every year.
- Clustering is highly application dependent and to some extent subjective.

Dimensionality reduction techniques

- Another unsupervised learning setting e.g PCA
- Dimensionality reduction is simply, the process of reducing the dimension of your feature set.
- Your feature set could be a dataset with a hundred columns (i.e features).
 - High-Dimensions = Lot of Features

Document classification

Features per document =
thousands of words/unigrams
millions of bigrams, contextual
information



Dimensionality reduction techniques

Dimensionality reduction

- Represent the data $x_1, \dots, x_n \in \mathbb{R}^d$ in a subspace of lower dimension with as little loss of information as possible
- **Advantages:**
 - Visualization
 - Lower computation and time complexity
 - Avoid overfitting and reduce noise

Dimensionality reduction techniques

- Problem settings:

- Given $x_1, \dots, x_n \in \mathbb{R}^d$
- find a k -dimensional subspace
- such that the data projected onto that space
- is as close to the original data as possible

- Common techniques:

- **Principal Component Analysis(PCA)**
- **Factor Analysis**
- **Linear Discriminant Analysis(LDA)**




please read FA and LDA

Dimensionality reduction techniques

- The most common and well known dimensionality reduction (PCA, LDA, FA)
 - **Principal Component Analysis(PCA):** PCA rotates and projects data along the direction of increasing variance. The features with the maximum variance are the principal components.


```
1: function PCA(parameter  $k$ , inputs  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ )
2:   compute sample mean  $\hat{\mu} := \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$  and center each input  $\mathbf{x}_i := \mathbf{x}_i - \hat{\mu}$ 
3:   compute scatter matrix  $S_n := XX^T$ , where  $X := (\mathbf{x}_1, \dots, \mathbf{x}_n)$ 
4:   compute eigenvectors  $\mathbf{w}_1, \dots, \mathbf{w}_k$  of  $S_n$ 
5:   (e.g eigenvalue, eigenvector= eig( $S_n$ ) gives  $\mathbf{w} = \mathbf{w}_1, \dots, \mathbf{w}_k$ )
6:   compute projected inputs  $\tilde{\mathbf{x}}_i := (\mathbf{w}_1^T \mathbf{x}_i, \dots, \mathbf{w}_k^T \mathbf{x}_i)^T, i = 1, \dots, n$ 
7:   return projected inputs  $\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_n$ 
8: end function
```

```
from numpy import array
from numpy import mean
from numpy import cov
from numpy.linalg import eig
# define a matrix
A = array([[1, 2], [3, 4], [5, 6]])
print(A)
# calculate the mean of each column
M = mean(A.T, axis=1)
print(M)
# center columns by subtracting column means
C = A - M
print(C)
# calculate covariance matrix of centered matrix
V = cov(C.T)
print(V)
# eigendecomposition of covariance matrix
values, vectors = eig(V)
print(vectors)
print(values)
# project data
P = vectors.T.dot(C.T)
print(P.T)
```

$$\text{Distance between clusters} = \frac{d_{1,3} + d_{1,4} + d_{1,5} + d_{2,3} + d_{2,4} + d_{2,5}}{6}$$

Dimensionality Reduction & Principal Component Analysis



- 1: **function** PCA(parameter k , inputs $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$)
- 2: compute sample mean $\hat{\mu} := \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$ and center each input $\mathbf{x}_i := \mathbf{x}_i - \hat{\mu}$
- 3: compute scatter matrix $S_n := XX^T$, where $X := (\mathbf{x}_1, \dots, \mathbf{x}_n)$
- 4: compute eigenvectors $\mathbf{w}_1, \dots, \mathbf{w}_k$ of S_n
- 5: (e.g., in MATLAB: [foo,W] = eig(S_n) gives $W = (\mathbf{w}_1, \dots, \mathbf{w}_k)$)
- 6: compute projected inputs $\tilde{\mathbf{x}}_i := (\mathbf{w}_1^T \mathbf{x}_i, \dots, \mathbf{w}_k^T \mathbf{x}_i)^T, i = 1, \dots, n$
- 7: **return** projected inputs $\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_n$
- 8: **end function**

