

# **Chapter-7**

## **Other Topics-Overview**

# Topics

---

- Ensemble learning vs
- Hybrid learning vs
- End-to-end learning
- Reinforcement learning

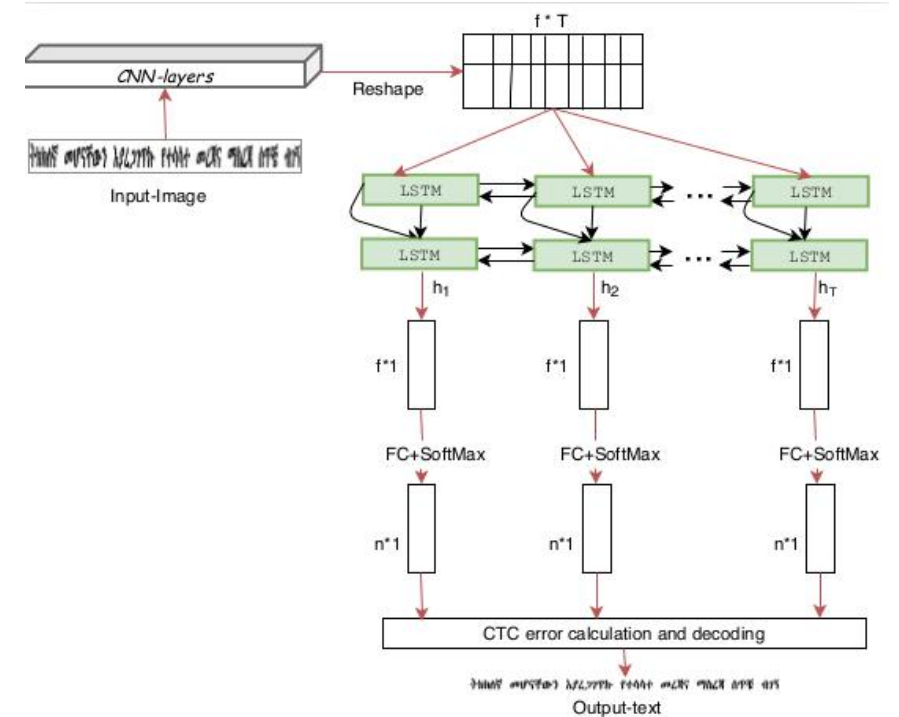
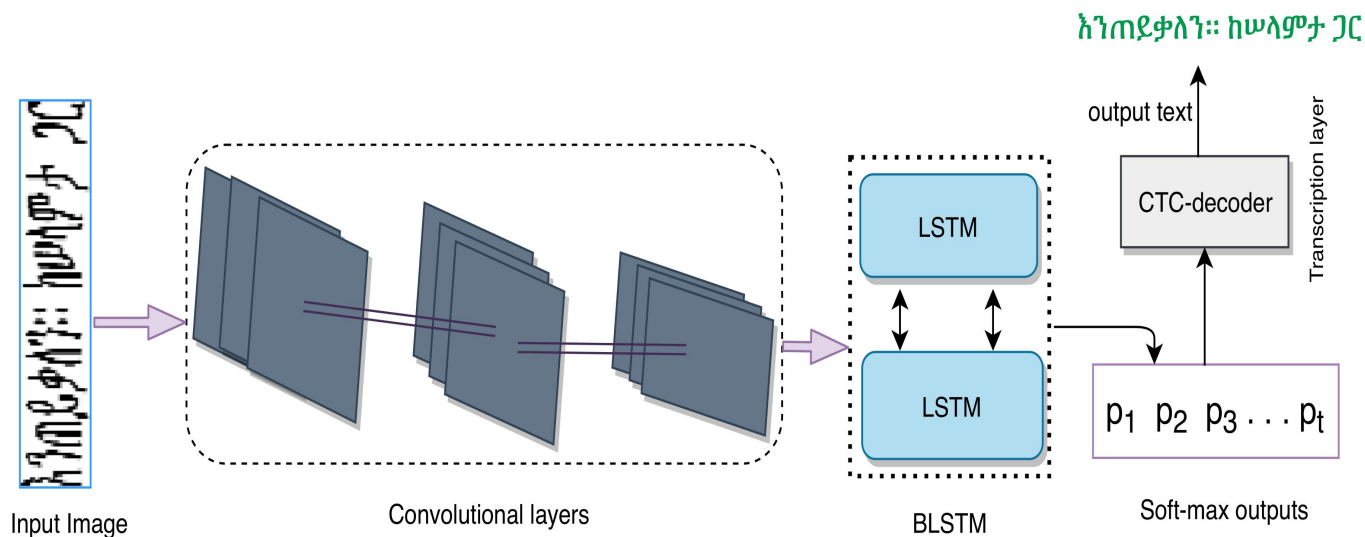
# Ensemble learning

---

- **Ensemble learning** is the process by which **multiple models**, such as classifiers or experts, are strategically generated and combined to solve a particular computational intelligence problem and to **improve model performance**.
  - **e.g AdaBoost**: algorithm for classification problems that add new machine learning models in a series where subsequent models attempt to fix the prediction errors made by prior models.
  - **Bagging**: Give equal weightage to all classifiers then majority voting
  - **Boosting**: Give weightage according to the accuracy of the classifier
- e.g a patient with a set of symptoms and taking opinion doctors
  - **Random Forest**: A subset of input features is chosen to form each training set

# Hybrid learning

- Combining two or more different machine learning techniques so as to build a model for solving a task.
  - e.g** using one unsupervised learner (or cluster) to pre-process the training data and one supervised learner (or classifier) to learn.
  - Using one as feature extractor while the other as a sequence learner



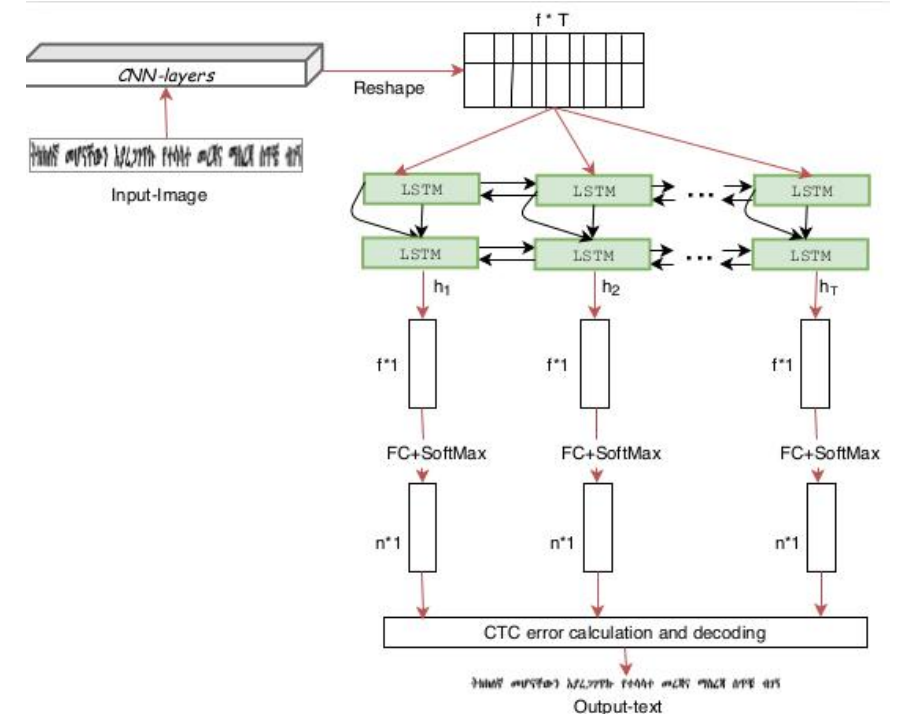
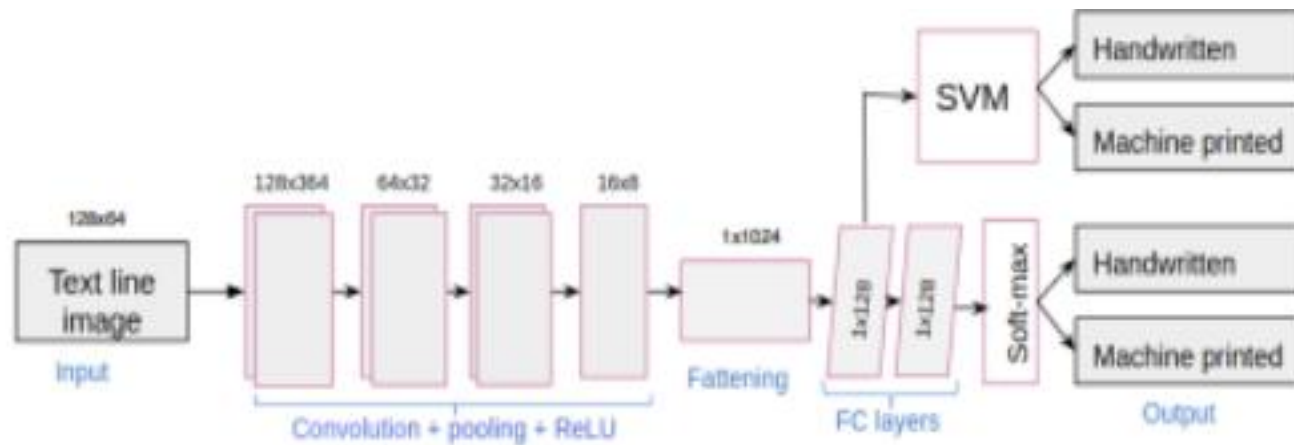
# Are they the same?

---

- **Ensemble** methods work independently to vote on an outcome while **hybrid** methods work together to predict one single outcome, which no voting element present in it.

# End-to- end learning

- CNN+SVM (hybrid but not end-to-end)
- CNN+LSTM+CTC ( train end-to-end)



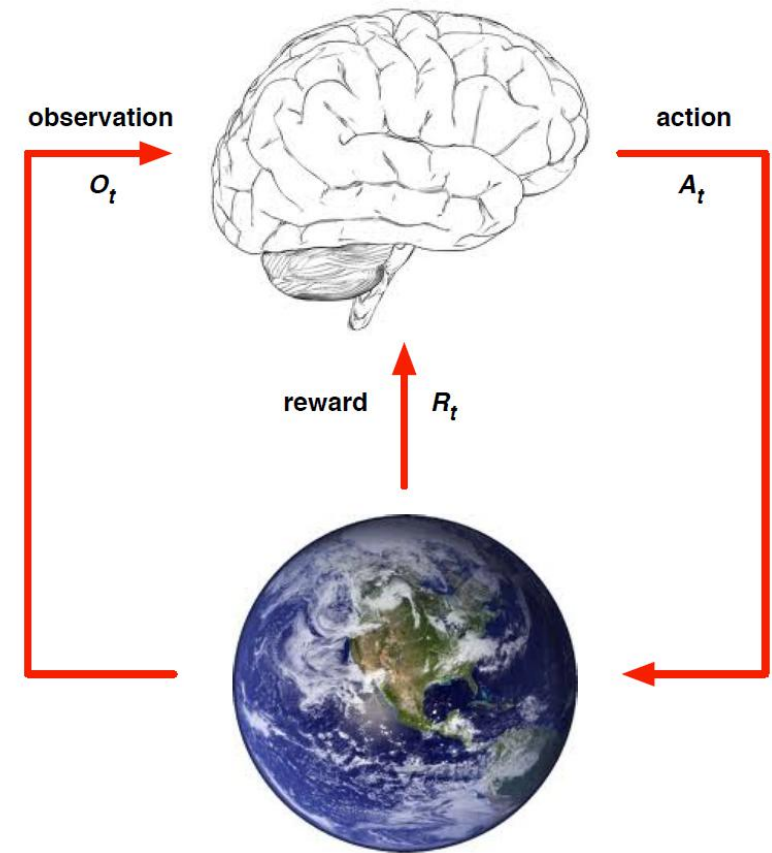
# Reinforcement learning (RL)

---

- Reinforcement learning (RL) is an area of ML concerned with how intelligent agents ought to **take actions** in an **environment** in order to maximize the notion of cumulative reward



# What is RL?

- Reward instead of data (no GT)
- Preferable states in the world are defined





# Supervised or Unsupervised? Both!




# Settings of RL

---

- **Agent:** An agent takes actions; for example, a drone making a delivery
- **Action (A):**  $A=\{a_1, a_2, \dots, a_m\}$  is the set of all possible moves the agent can make: e.g. running right or left, jumping high or low, buying or selling
- **Environment:** The world through which the agent moves, and which responds to the agent.
  - The environment takes the agent's current state and action as input, and returns as output the agent's reward and its next state.
  - e.g. If you are the agent, the environment could be the **laws of physics** and the **rules of society** that process your actions and determine the consequences of them.
- **State (S):** A state is an immediate situation in which the agent finds itself in relation to other significant things such as tools, obstacles, enemies or prizes
- **Reward (R):** A reward is the feedback by which we measure the success or failure of an agent's actions in a given state
- **Policy ( $\pi(s)$ ):** is the strategy that the agent employs to determine the next action based on the current state.
- **Value ( $V\pi(s)$ ):** is defined as the expected long-term return of the current state under policy  $\pi$
- **Model:**
  - predict next state  $P^a_{ss'}$
  - predict the next reward  $P^a_s$


# Example

- Reward: -1 per time step
- States: Agent's location
- Actions: N,E,S,W




# Policy

- Arrows represent policy  $\pi(s)$  for each state  $s$




# Value Function

- Numbers represent value  $V_{\pi}(s)$  for each state  $s$




# Model

- Grid layout represents transition model  $P_{ss'}^a$ ,
- Numbers represent immediate reward  $R_s^a$  from each state  $s$  (same for all  $a$ )






# Example...

- Given an image that represents a state, a convolutional net can rank the actions possible to perform in that state
- The below image illustrates what a policy agent does, mapping a state to the best action.  
 $a = \pi(s)$
- A policy maps a state to an action.



# Neural Network as Optimal Policy Estimator

[https://www.youtube.com/watch?time\\_continue=77&v=YOW8m2YGtRg](https://www.youtube.com/watch?time_continue=77&v=YOW8m2YGtRg)

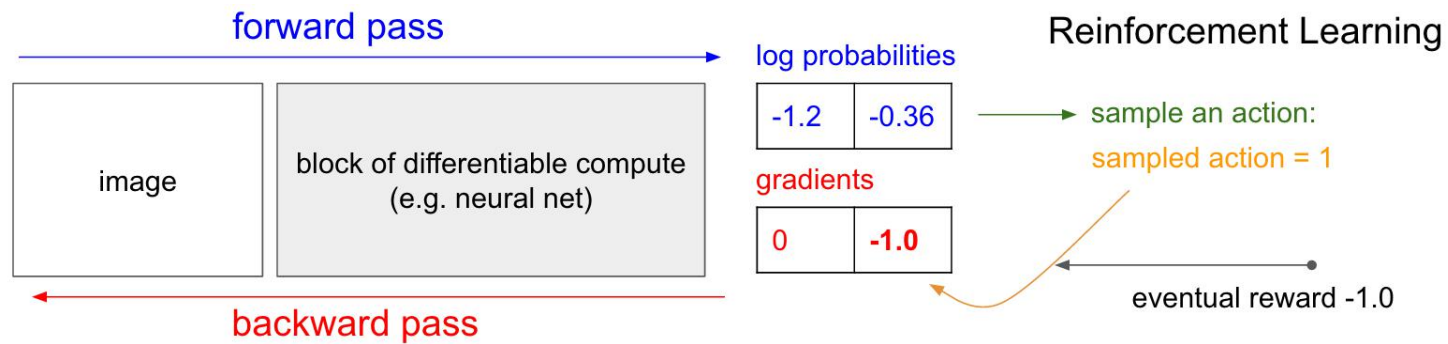
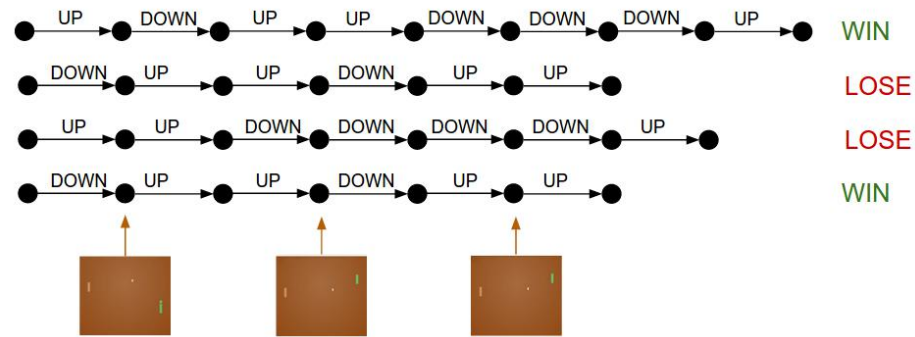


-->a pong game that simulates table tennis  
(210\*160\*3) input image



# How to Train such a Network ?

- “generate a dataset” to then perform supervised learning on-the-fly



# Characteristics of RL

---

- **Input:** The input should be an initial state from which the model will start
- **Output:** There are many possible output as there are variety of solution to a particular problem
- **Training:** The training is based upon the input, The model will return a state and the user will decide to reward or punish the model based on its output.
- The model keeps continues to learn.
- The best solution is decided based on the maximum reward.
- There is no supervisor, only a reward signal
- Feedback is delayed, not instantaneous
- Agent's actions affect the subsequent data it receives

# Applications of RL

---

- Robotics for Industrial Operations
- Supply Chain & Logistics
- Traffic Control
- Bidding & Advertising
- Recommender Systems
- Load Balancing

# Further Readings

---

- Deep Q-Learning
- Semi-supervised SVM (S3VM)
- Policy vs value vs policy+value-based learning
- Bandit problems and online learning
- Sources for further reading on RL
  - [Udacity (Georgia Tech.)] CS7642 Reinforcement Learning
  - [Stanford] CS229 Machine Learning - Lecture 16: Reinforcement Learning by Andrew Ng
  - [UC Berkeley] Deep RL Bootcamp
  - [UC Berkeley] CS294 Deep Reinforcement Learning by John Schulman and Pieter Abbeel
  - [CMU] 10703: Deep Reinforcement Learning and Control, Spring 2017
  - [MIT] 6.S094: Deep Learning for Self-Driving Cars
  - Lecture 2: Deep Reinforcement Learning for Motion Planning

or just go to this link <https://wiki.pathmind.com/deep-reinforcement-learning>

# Information

---

Chapter 1: 8%

Chapter 2: 25%

chapter 3: 10%

chapter 4: 25%

chapter 5: 5%

chapter 6: 15%

chapter 7: 2%

General= 10%

5-10 questions... you have already done 20% of the questions

Exam-date: TBA

# Tips

---

Demo on overleaf

- ^ <https://www.overleaf.com/project>

sample exam questions