

# Introduction to Programming with C++

Data Types

# Objective

- In this chapter we'll discuss:
  - Data types in C++
  - How you declare and initialize variables
  - What literals are and how you define them
  - Binary and hexadecimal integers
  - How calculations work
  - How you can fix the value of a variable
  - How to create variables that store characters
  - What the auto keyword does
  - What lvalues and rvalues are

# Variables, Data, & Data Types

- A variable is a named piece of memory that you define.
- Each variable only stores data of a particular type.
- Every variable has a type that defines the kind of data it can store.
- Each fundamental type is identified by a unique type name that is a keyword.
- The compiler makes extensive checks to ensure that you use the right data type in any given context.
- It will also ensures compatibility during operations
- The compiler detects and reports attempts to combine data of different types that are incompatible.

# Variables, Data, & Data Types

- Numerical values fall into two broad categories:
  - Integers-whole numbers in other words,
  - Floating-point values, which can be non-integral.
- There are several fundamental C++ types in each category,
- Each of these can store a specific range of values.

# Defining Integer Variables

- Here's a statement that defines an integer variable:

```
int appleCount;
```

- This defines a variable of type `int` with the name `appleCount`.
- The variable will contain some arbitrary junk-value.
- You can and should specify an initial value when you define the variable, like this:

```
int appleCount {15}; // Number of apples
```

- The initial value for `appleCount` appears between the braces following the name so it has the value 15.
- The braces enclosing the initial value is called an initializer list.

# Defining Integer Variables

- You'll meet situations later in the course where an initializer list will have several values between the braces.
- You don't have to initialize variables when you define them but it's a good idea to do so.
- Ensuring variables start out with known values makes it easier to work out what is wrong when the code doesn't work as you expect.
- Type `int` is typically 4 bytes, which can store integers from **-2,147,483,648** to **+2,147,483,647**.
- This covers most situations, which is why `int` is the integer type that is used most frequently.

# Defining Integer Variables

- Here are definitions for three variables of type `int`:

```
int appleCount {15}; // Number of apples
int orangeCount {5}; // Number of oranges
int totalFruit {appleCount + orangeCount}; // Total number of fruit
```

- The initial value for `totalFruit` is the sum of the values of two variables defined previously.
- This demonstrates that the initial value for a variable can be an expression.

The statements that define the two variables in the expression for the initial value for `totalFruit` must appear earlier in the source file. Otherwise the definition for `totalFruit` won't compile.

# Defining Integer Variables

- The initial value between the braces should be of the same type as the variable you are defining.
- If it isn't, the compiler will have to convert it to the required type.
- If the conversion is to a more limited range of values, it potentially could lead to loss of information
- Hence, the compiler won't convert the value but just flag it as an error.
- An example would be if you specified the initial value for an integer variable that is not an integer—1.5 for example.
- You might do this by accident when entering the value 15 for `appleCount`.



# Defining Integer Variables

- A conversion to a type with a more limited range of values is called a narrowing conversion.
- There are two other ways for initializing a variable.
- Functional notation looks like this:

```
int orangeCount(5);
```

```
int totalFruit(appleCount + orangeCount);
```

- Alternatively you could write this:

```
int orangeCount = 5;
```

```
int totalFruit = appleCount + orangeCount;
```

# Defining Integer Variables

- While both of these possibilities are valid
- It's recommended that you adopt the initializer list form.
- This is the most recent syntax that was introduced in C++ 11 to standardize initialization.
- It is preferred because it enables you to initialize just about everything in the same way.
- There's one exception that we'll discuss later
- You can define and initialize more than one variable of a given type in a single statement.
- For example:

```
int footCount {2}, toeCount {10}, headCount {1};
```

- Most of the time it's better to define each variable in a separate statement.
- This makes the code more readable and you can explain the purpose of each of them in a comment.

# Signed Integer Types

Type Name	Typical Size (bytes)	Range of Values
<code>signed char</code>	1	-128 to 127
<code>short</code> <code>short int</code>	2	-256 to 255
<code>int</code>	4	4 -2,147,483,648 to +2,147,483,647
<code>long</code> <code>long int</code>	4	4 -2,147,483,648 to +2,147,483,647
<code>long long</code> <code>long long int</code>	8	8 -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807

- Type `signed char` is typically 1 byte;
- The number of bytes occupied by the others depends on the compiler.
- Where two type names appear in the left column, the abbreviated name that comes first is commonly used
- So you will usually see `long` used rather than `long int`.
- Each type will have at least as much memory as the one that precedes it in the list.

# Unsigned Integer Types

- There are circumstances where you don't need to store negative numbers.
- The number of students in a class or the number of parts in an assembly are always positive integers.
- You can specify integer types that only store non-negative values by prefixing any of the names of the signed integer types with the `unsigned` keyword
- For example
  - `unsigned char`
  - `unsigned short`
  - `unsigned long`

# Unsigned Integer Types

- Each unsigned type is a different type from the signed type but occupies the same amount of memory.
- Type char is a different integer type from both signed char and unsigned char.
- Type char stores a character code and can be a signed or unsigned type depending on your compiler.
- If the constant `CHAR_MIN` in the `climits` header is 0, then char is an unsigned type with your compiler.

# Unsigned Integer Types

- Here are some examples of variables of some of these types:

```
signed char ch {20};  
long temperature {-50L};  
long width {500L};  
long long height {250LL};  
unsigned int toe_count {10U};  
unsigned long angel_count {1000000UL};
```

- Note how you write constants of type `long` and type `long long`.
- You must append L to the first and LL to the second.
- If there is no suffix, an integer constant is of type `int`.
- You can use lowercase for the L and LL suffixes but it can easily be confused with the digit 1.
- Unsigned integer constants have u or U appended.

# Fixed Value Variables

- Sometimes you'll want to define variables with values that are fixed and must not be changed.
- You use the `const` keyword in the definition of a variable that must not be changed.
- For example:

```
const unsigned int toe_count {2U};
```

- The `const` keyword tells the compiler that the value of `toe_count` must not be changed.
- A statement that attempts to modify the value of `toe_count` will be flagged as an error during compilation.
- You can use the `const` keyword to fix the value of variables of any type.

# Integer Literals

- Constants of any kind, such as 42, or 2.71828, 'z', or "Mark Twain", are referred to as literals.
- Every literal will be of some type.
- For example,
  - 42 is an integer literal,
  - 2.71828 is floating-point literal,
  - 'z' is a character literal
  - "Mark Twain" is a string literal.



# Integer Literals

- You can write integer literals in a very straightforward way.
- Here are some examples of decimal integers:

`-123L +123 123 22333 98U -1234LL 12345ULL`

- You have seen that unsigned integer literals have `u` or `U` appended.
- Literals of types long and type long long have `L` or `LL` appended respectively
- If they are unsigned, they also have `u` or `U` appended.
- The `U` and `L` or `LL` can be in either sequence.
- You could omit the `+` in the second example, as it's implied by default
- But if you think putting it in makes things clearer, that's not a problem.

# Integer Literals

- The literal `+123` is the same as `123` and is of type `int` because there is no suffix.
- The fourth example is the number that you would normally write as 22,333, but you must not use commas in an integer literal.
- Here are some statements using some of these literals:

```
unsigned long age {99UL}; // 99uL would be OK too
unsigned short {10u}; // There is no specific
literal type for short
long long distance {1234567LL};
```

- You can't write just any old integer value as an initial value for a variable.
- An initializing value must be within the permitted range for the type of variable as well as match the type.
- A literal in an expression must be within the range of some type.

# Hexadecimal Literals

- You can write integer literals as hexadecimal values.
- You prefix a hexadecimal literal with `0x` or `0X`
- `0x999` is a hexadecimal number of type `int` with three hexadecimal digits.
- Plain old `999`, on the other hand, is a decimal value of type `int` with decimal digits
- So the value will be completely different.
- Here are some more examples of hexadecimal literals:

Hexadecimal literals:	<code>0x1AF</code>	<code>0x123U</code>	<code>0xAL</code>	<code>0xcad</code>	<code>0xFF</code>
Decimal literals:	<code>431</code>	<code>291U</code>	<code>10L</code>	<code>3245</code>	<code>255</code>

# Hexadecimal Literals

- A major use for hexadecimal literals is to define particular patterns of bits.
- Each hexadecimal digit corresponds to 4 bits so it's easy to express a pattern of bits as a hexadecimal literal.
- The red, blue, and green components (RGB values) of a pixel color are often expressed as three bytes packed into a 32-bit word.
- Here are some examples:

```
unsigned int color {0x0f0d0eU}; // decimal 986,382
```

```
int mask {0xFF00FF00}; // Four bytes specified as FF, 00, FF, 00
```

```
unsigned long value {0xdeadLU}; // decimal 57,005
```

# Octal Literals

- You can write integer literals as octal value.
- You identify a number as octal by writing it with a leading zero.
- You can also write a binary integer literal as a sequence of binary digits (0 or 1) prefixed by 0b or 0B.

**Octal literals:**     657             0443U             012L             6255             377

**Decimal literals:**     431             291U             10L             3245             255

**Binary Literals:**     0B110101111     0b100100011U     0b1010L     0B11001101     0b11111111

- Binary literals were introduced by the C++ 14 standard
- As far as your compiler is concerned, it doesn't matter which number base you choose when you write an integer value.
- Ultimately it will be stored as a binary number.
- The different ways for writing an integer are there just for your convenience.
- You choose one or other of the possible representations to suit the context.

# Calculations with Integers

- To begin with, let's get some bits of terminology out of the way.
- An operation such as addition or multiplication is defined by an operator
- The values that an operator acts upon are called operands
- For example, in an expression such as  $2*3$ , the operands are 2 and 3.
- Operators such as multiplication that require two operands are called binary operators.
- Operators that require one operand are called unary operators.
- An example of a unary operator is the minus sign
- For example in the expression `-width`, the minus sign negates the value of width

# Calculations with Integers

Operator	Operation	Example
+	Addition	4+5 is 9
-	Subtraction	4-5 is -1
*	Multiplication	4*5 is 20
/	Division	6/4 is 1 (integer devision)
%	Modulus	6%4 is 2 (remainder of the devision)
=	Assignment	<code>long x = 57;</code> // the variable x will assume the value of 57 from this point on

# More on Assignment Operations

- You can assign a value to more than one variable in a single statement.
- For example:

```
int a {}, b {}, c {5}, d{4};  
a = b = c*c - d*d;
```

- The second statement calculates the value of the expression  $c*c - d*d$  and stores the result in b, so b will be set to 9.
- Next the value of b is stored in a so a will also be set to 9.
- You can have as many repeated assignments like this as you want.
- The operand on the left of an assignment can be a variable or an expression



# More on Assignment Operations

- If it is an expression, the result must be an lvalue.
- An lvalue represents a persistent memory location so a variable is an lvalue.
- Every expression in C++ results in either an lvalue or an rvalue.
- An rvalue is a result that is not an lvalue, so it is transient.
- The result of the expression  $c * c - d * d$  in the statement above is an rvalue.
- The compiler allocates a temporary memory location to store the result of the expression
- Once the statement has been executed, the result and the memory it occupies is discarded.

- The difference between lvalues and rvalues is not important now
- But it will become very important when you delve into functions and classes.

# The op= Assignment Operators

- In Ex2\_01.cpp, there was a statement that you could write more economically:

```
feet = feet % feetPerYard;
```

- This statement could be written using an op= assignment operator.
- The op= assignment operators are so called because they're composed of an operator and an assignment operator =.
- You could write the previous statement as:

```
feet %= feetPerYard;
```

- This is exactly the same operation as the previous statement.
- In general, an op= assignment is of the form:

```
lhs op= rhs;
```

# The op= Assignment Operators

- `lhs` represents a variable of some kind that is the destination for the result of the operator.
- `rhs` is any expression.
- This is equivalent to the statement:

```
lhs = lhs op (rhs);
```

- The parentheses are important because you can write statements such as:

```
x *= y + 1;
```

- This is equivalent to:

```
x = x*(y + 1);
```

- Without the implied parentheses, the value stored in `x` would be the result of `x*y+1`, which is quite different.
- You can use a range of operators for `op` in the `op=` form of assignment.

# The op= Assignment Operators

Operation	Operator	Operation	Operator
Addition	+=	Bitwise AND	&=
Subtraction	-=	Bitwise OR	=
Multiplication	*=	Bitwise exclusive OR	^=
Division	/=	Shift left	<<=
Modulus	%=	Shift right	=>>

- Note that there can be no spaces between op and the =.
- If you include a space, it will be flagged as an error.
- You can use += when you want to increment a variable by some amount.
- For example, the following two statements have the same effect:

# using Declarations & Directives

- There were a lot of occurrences of `std::cin` and `std::cout` in `Ex2_01.cpp`.
- You can eliminate the need to qualify a name with the namespace name in a source file with a using declaration.
- Here's an example:

```
using std::cout;
```

- This tells the compiler that when you write `cout`, it should be interpreted as `std::cout`.
- This saves typing and makes the code look a little less cluttered.

# using Declarations & Directives

- You can apply using declarations to names from any namespace, not just `std`.
- A using directive imports all the names from a namespace.
- Here's how you could use any name from the `std` namespace without the need to qualify it:

```
using namespace std; // Make all the names in std
available without qualification
```

- With this at the beginning of a source file, you don't have to qualify any name that is defined in the `std` namespace.

- At first sight this seems an attractive idea.
- The problem is it defeats a major reason for having namespaces.
- It is unlikely that you know all the names that are defined in `std` and with this using directive you have increased the probability of accidentally using a name from `std`.

# The sizeof Operator

- You use the sizeof operator to obtain the number of bytes occupied by a type, or by a variable, or by the result of an expression.
- Here are some examples of its use:

```
int height {74};  
std::cout << "The height variable occupies "  
           << sizeof height << " bytes."  
           << std::endl;  
std::cout << "Type \"long long\" occupies " <<  
           sizeof (long long) << " bytes." <<  
           std::endl;  
std::cout << "The expression height * height/2"  
           << occupies " << sizeof (height*height/2)  
           << " bytes." << std::endl;
```

# The sizeof Operator

- These statements show how you can output the size of a variable, the size of a type, and the size of the result of an expression.
- To use sizeof to obtain the memory occupied by a type, the type name must be between parentheses.
- You also need parentheses around an expression with `sizeof`.
- You don't need parentheses around a variable name, but there's no harm in putting them in.
- The result that sizeof produces is of type `size_t`.
- `size_t` is an unsigned integer type that is defined in the Standard Library header `cstdint.h`.



# Incrementing and Decrementing Integers

- You can use += and -= with the literal 1 on the rvalue to increment and decrement the lvalue by one
- There are two other operators that can perform the same tasks.
- They're called the increment operator and the decrement operators, ++ and -- respectively.
- These are unary operators that you can apply to an integer variable.
- The following statements that modify count have exactly the same effect:

```
int count {5};  
count = count + 1;  
count += 1;  
++count;
```

# Incrementing and Decrementing Integers

- Each statement increments count by 1.
- Using the increment operator is clearly the most concise.
- The effect in an expression is to increment the value of the variable and then to use the incremented value in the expression.
- For example, suppose count has the value 5, and you execute this statement:

```
total = ++count + 6;
```

- The increment and decrement operators execute before any other binary arithmetic operators in an expression.

# Postfix Increment and Decrement Operations

- The postfix form of ++ increments the variable to which it applies after its value is used in context.
- For example, you can rewrite the earlier example as follows:

```
total = count++ + 6;
```

- With an initial value of 5 for count, total is assigned the value 11.
- count will then be incremented to 6.
- The preceding statement is equivalent to the following statements:

```
total = count + 6;  
++count;
```

- These rules also apply to the decrement operator.

# Postfix Increment and Decrement Operations

- You must not apply the prefix form of these operators to a given variable more than once in an expression.

- Suppose count has the value 5, and you write this:

```
total = ++count * 3 + ++count * 5;
```

- Because the statement modifies the value of count more than once, the result is undefined.
- You should get an error message from the compiler with this statement.
- Note also that the effects of statements such as the following are undefined:

```
k = ++k + 1;
```

- Although such expressions are undefined according to the C++ standard, this doesn't mean that your compiler won't compile them.
- It just means that there is no guarantee of consistency in the results.

# Defining Floating-Point Variables

- You use floating-point variables whenever you want to work with values that are not integral.
- There are three floating-point data types.
  - `float` - Single precision floating-point values
  - `double` - Double precision floating-point values
  - `long double` - Double-extended precision floating-point values
- The term “precision” refers to the number of significant digits in the mantissa.
- The types are in order of increasing precision

# Defining Floating-Point Variables

- Note that the precision only determines the number of digits in the mantissa.
- The range of numbers that can be represented by a particular type is determined by the range of possible exponents.
- The precision and range of values aren't prescribed by the C++ standard so what you get with each type depends on your compiler.
- This will depend on what kind of processor is used by your computer and the floating-point representation it uses.
- Type long double will provide a precision that's no less than that of type double, and type double will provide a precision that is no less than that of type float.

# Defining Floating-Point Variables

- Typically,
  - `float` provides 7 digits precision
  - `double` provides 15 digits,
  - `long double` provides 19 digits precision
  - `double` and `long double` have the same precision with some compilers.
- Typical ranges of values on an Intel processor are

Type	Precision	Range
<code>float</code>	7	<b><math>1.2 \times 10^{-38}</math> to <math>3.4 \times 10^{38}</math></b>
<code>double</code>	15	<b><math>2.2 \times 10^{-308}</math> to <math>1.8 \times 10^{308}</math></b>
<code>long double</code>	19	<b><math>3.3 \times 10^{-4932}</math> to <math>1.2 \times 10^{4932}</math></b>

# Floating-Point Literals

- You can see from the code fragment in the previous section that float literals have f (or F) appended and long double literals have L (or l) appended.
- Floating point literals without a suffix are of type double.
- A floating-point literal includes either a decimal point, or an exponent, or both
- A numeric literal with neither is an integer.
- An exponent is optional in a floating-point literal
- It represents a power of 10 that multiplies the value.
- An exponent must be prefixed with e or E and follows the value.
- Here are some floating-point literals that include an exponent:

`5E3` (5000.0) `100.5E2` (10050.0) `2.5e-3` (0.0025) `-0.1E-3L` (-0.0001L) `.345e1F` (3.45F)

- Exponents are particularly useful when you need to express very small or very large values.



# Floating-Point Calculations

- You write floating-point calculations in the same way as integer calculations.
- For example:

```
const double pi {3.1414926};  
double a {0.75}; // Thickness of a pizza  
double z {5.5}; // Radius of a pizza  
double volume {}; // Volume of pizza - to be calculated  
volume = pi*z*z*a;
```

- The modulus operator, %, can't be used with floating-point operands.
- You can also apply the prefix and postfix increment and decrement operators

# Reading Assignment

- Mathematical Functions (page 40)
- Formatting Stream Output(Page 43)